

---

# Bespoke Cloud Technology

# Introduction to Cloud and the Cloud Native Model

---

# What Is Cloud Computing?

- On-demand resources
- Elastic capacity
- Shared responsibility model
- API-driven infrastructure

# On-demand resources

What it means:

- Infrastructure resources (compute, storage, network, managed services) are:
  - provisioned **automatically**
  - available in **seconds or minutes**
  - created **without human interaction** from the provider side

# On-demand resources

Security / audit perspective:

- No manual approval does not mean no control
- Control shifts to:
  - IAM
  - policies
  - guardrails

# On-demand resources

Security / audit perspective:

- IAM roles and permissions
- Cloud audit logs (API calls)
- Infrastructure-as-Code repositories  
(Terraform, CloudFormation)

# Elastic capacity

What it means:

- Systems can:
  - scale up and down automatically
  - react dynamically to load
  - avoid fixed, static capacity

# Elastic capacity

Security implications:

- Risk assessment changes:
  - not how many servers exist
  - but how scaling is controlled
- Common risks:
  - uncontrolled cost growth
  - amplification of DoS attacks
  - missing quotas or limits

# Elastic capacity

Typical audit evidence:

- Autoscaling configurations
- Resource quotas and limits
- Budget alerts and monitoring rules
- Kubernetes HPA / cloud autoscaling policies

# Shared responsibility model

What it means:

- Security responsibilities are shared between:
  - the cloud provider
  - the customer
- The exact boundary depends on:
  - IaaS vs PaaS vs SaaS
  - the specific service used

# Shared responsibility model

Security / audit perspective:

- One of the most common audit failures:  
“We assumed the provider was responsible for this.”
- Audits must clearly define:
  - what is out of scope
  - what remains customer responsibility

# Shared responsibility model

- Typical audit evidence:
  - Responsibility matrices
  - Provider security documentation
  - Architecture and data-flow diagrams

# API-driven infrastructure

What it means:

- Everything in the cloud is:
  - created via APIs
  - managed via APIs
  - logged via APIs

# API-driven infrastructure

## Security implications:

- IAM becomes the primary security perimeter
- API abuse is a real attack vector
- Service accounts, tokens, and workload identities are critical assets

# API-driven infrastructure

Audit value (very important):

- Full traceability:
  - who did what
  - when
  - from where
- Enables:
  - incident reconstruction
  - forensic analysis
  - policy enforcement validation

# API-driven infrastructure

Typical audit evidence:

- Cloud audit logs
- Service account usage logs
- API call history
- GitOps and CI/CD pipeline logs

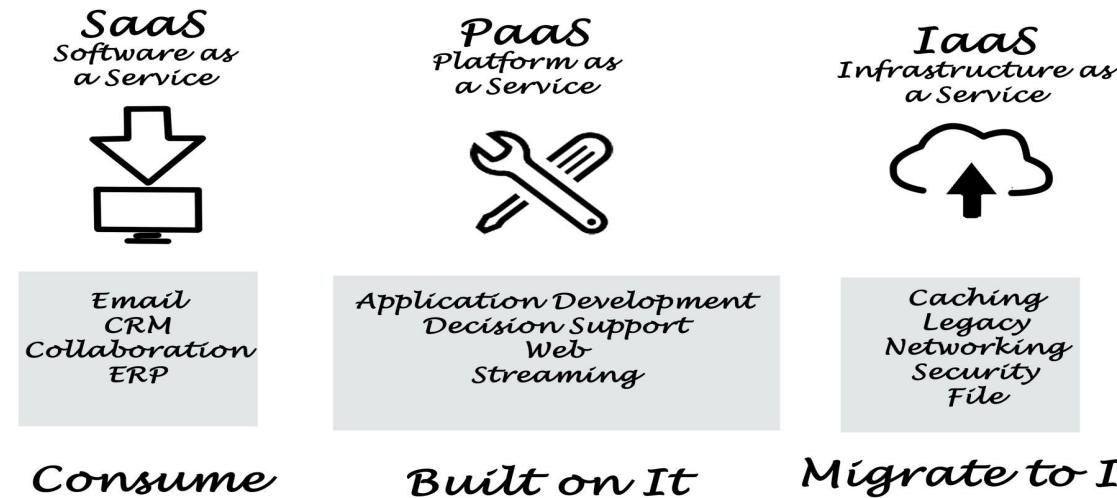
# Cloud Provider Tools for Audit Evidence

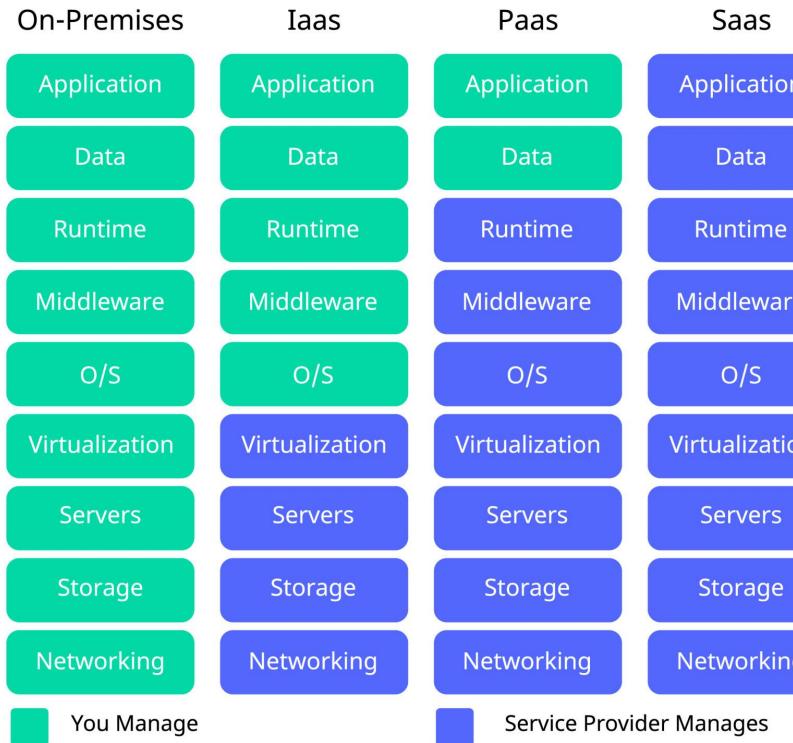
Key Areas and Native Tools for AWS, Google Cloud, and Azure

| Audit Evidence Category   | What Auditors Verify         | AWS   | Google Cloud  | Azure  |
|---|------------------------------|---|---|--|
|  Identity & Access (WHO)             | Roles, Least Privilege       | IAM, Access Analyzer  | Cloud IAM, Org Policies   | Azure AD, RBAC   |
|  API Activity (WHAT / WHEN)          | API Call Logs, Admin Actions |  CloudTrail              |  Cloud Audit Logs        |  Activity Logs      |
|  Configuration State (HOW IT IS SET) | Drift, Compliance            |  AWS Config              |  Config Validator        |  Azure Policy       |
|  Change Control (HOW IT CHANGED)     | IaC, Versioned Changes       |  CloudFormation          |  Deployment Manager      |  ARM / Bicep        |
|  Security Posture (IS IT SAFE)       | Threats, Misconfigurations   |  Security Hub, GuardDuty |  Security Command Center |  Defender for Cloud |
|  Usage & Limits (HOW MUCH)           | Quotas, Budgets              |  Service Quotas, Budgets |  Quotas, Budget Alerts   |  Quotas, Cost Mgmt  |
|  Logs & Forensics (WHAT HAPPENED)    | Log Data, Retention          |  CloudWatch Logs         |  Cloud Logging           |  Log Analytics      |

# IaaS vs PaaS vs SaaS

- Different responsibility models
- Different audit controls



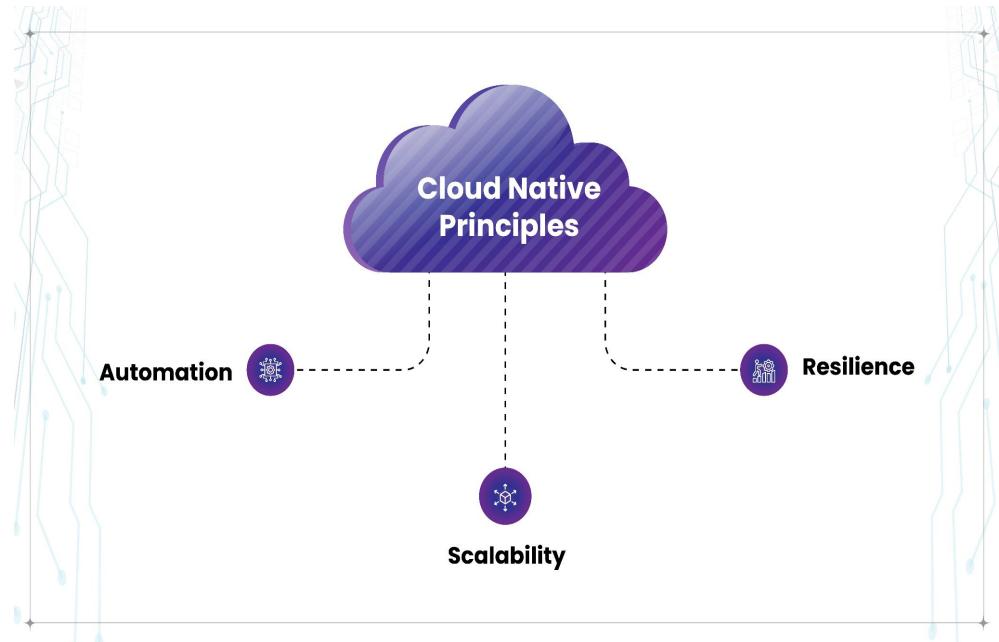


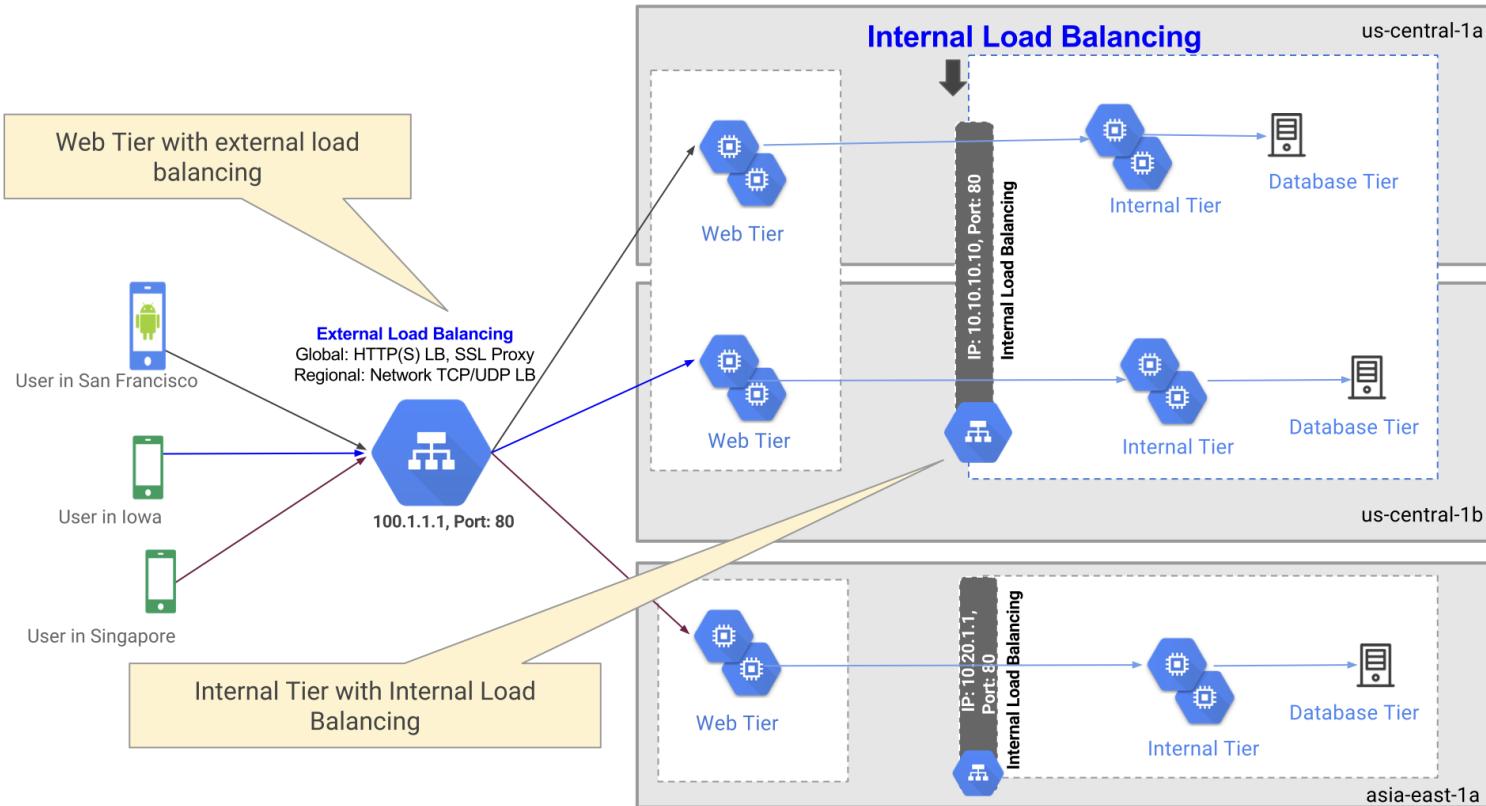
# Cloud-Native Characteristics

- Immutable infrastructure
- Automation by default
- Horizontal scaling
- API-first design

# Scalability, Resilience, Automation

- Failure is expected
- Self-healing systems
- Scaling without human approval





# Major Cloud Providers



Virtual Servers  
Platform-as-a-Service  
Serverless Computing  
Docker Management  
Kubernetes Management  
Object Storage  
Archive Storage  
File Storage  
Global Content Delivery  
Managed Data Warehouse



Instances  
Elastic Beanstalk  
Lambda  
ECS  
EKS  
S3  
Glacier  
EFS  
CloudFront  
Redshift



VM Instances  
App Engine  
Cloud Functions  
Container Engine  
Kubernetes Engine  
Cloud Storage  
Coldline  
ZFS / Avere  
Cloud CDN  
Big Query



VMs  
Cloud Services  
Azure Functions  
Container Service  
Kubernetes Service  
Block Blob  
Archive Storage  
Azure Files  
Delivery Network  
SQL Warehouse

# CLOUD COMPARISON

## AZURE vs. AWS vs. GOOGLE COMPUTE

|  Azure |  aws                      |   |  |
|---|--|--|--|
| Available Regions   | Azure Regions  | AWS Regions and Zones  | Google Compute Regions & Zones   |
| Compute Services  |  Virtual Machines           |  Elastic Compute Cloud (EC2)              |  Compute Engine                             |
| App Hosting   |  Azure Cloud Services       |  Amazon Elastic Beanstalk                 |  Google App Engine                          |
| Serverless Computing  |  Azure Functions            |  AWS Lambda                               |  Google Cloud Functions                     |
| Container Support   |  Azure Container Service    |  EC2 Container Service                    |  Container Engine                           |
| Scaling Options   |  Azure Autoscale            |  Auto Scaling                             |  Autoscaler                                 |
| Object Storage  |  Azure Blob Storage         |  Amazon Simple Storage (S3)               |  Cloud Storage                              |
| Block Storage   |  Azure Managed Storage      |  Amazon Elastic Block Storage             |  Persistent Disk                            |
| Content Delivery Network (CDN)  |  Azure CDN                  |  Amazon CloudFront                        |  Cloud CDN                                  |
| SQL Database Options  |  Azure SQL Database         |  Amazon RDS                               |  Cloud SQL                                  |
| NoSQL Database Options  |  Azure DocumentDB           |  AWS DynamoDB                             |  Cloud Datastore                            |
| Virtual Network   |  Azure Virtual Network      |  Amazon VPC                               |  Cloud Virtual Network                      |
| Private Connectivity  |  Azure Express Route        |  AWS Direct Connect                       |  Cloud Interconnect                         |
| DNS Services  |  Azure Traffic Manager      |  Amazon Route 53                          |  Cloud DNS                                  |
| Log Monitoring  |  Azure Operational Insights |  Amazon CloudTrail                        |  Cloud Logging                              |
| Performance Monitoring  |  Azure Application Insights |  Amazon CloudWatch                        |  Stackdriver Monitoring                     |
| Administration and Security   |  Azure Active Directory     |  AWS Identity and Access Management (IAM) |  Cloud Identity and Access Management (IAM) |
| Compliance  |  Azure Trust Center         |  AWS CloudHSM                             |  Google Cloud Platform Security             |
| Analytics   |  Azure Stream Analytics     |  Amazon Kinesis                           |  Cloud Dataflow                             |
| Automation  |  Azure Automation           |  AWS Opsworks                             |  Compute Engine Management                  |
| Management Services & Options   |  Azure Resource Manager     |  Amazon CloudFormation                    |  Cloud Deployment Manager                   |
| Notifications   |  Azure Notification Hub     |  Amazon Simple Notification Service (SNS) | None   |
| Load Balancing  |  Load Balancing for Azure   |  Elastic Load Balancing                   |  Cloud Load Balancing                       |

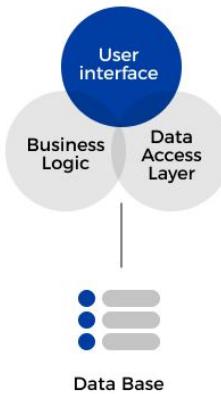
# Microservice Architecture vs Monolith

---

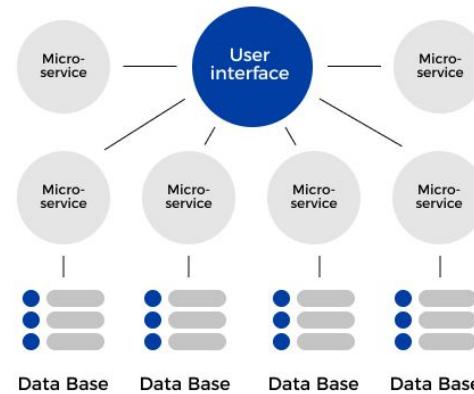
# Monolith vs Microservices

- Single unit vs distributed

**MONOLITHIC ARCHITECTURE**



**MICROSERVICE ARCHITECTURE**



# Microservice Design Principles

- Single responsibility
- Independent deployment
- Explicit interfaces
- Decentralized data

# Single Responsibility

What it means:

- Each microservice owns one business capability
- Changes in one domain should not force changes in others

Why it matters:

- Smaller codebases → easier reviews
- Smaller blast radius during failures or incidents

Security / audit perspective:

- Clear ownership of:
  - data
  - logic
  - access rights
- Easier threat modeling and data classification

# Independent Deployment

## What it means

- Each service:
  - has its own CI/CD pipeline
  - can be deployed independently
  - has its own lifecycle

## Why it matters

- No “big bang” deployments
- Faster incident response and rollback

# Independent Deployment

Security / audit perspective

- Every deployment:
  - is logged
  - has an author
  - has a timestamp
  - Strong change-management evidence

Trade-offs

- More pipelines to govern
- Higher operational complexity

# Explicit Interfaces

## What it means

- Services communicate only via:
  - APIs
  - events
  - No shared databases or memory

## Why it matters

- Clear contracts
- Easier versioning and compatibility control

# Explicit Interfaces

Security / audit perspective

- Natural enforcement points:
  - authentication
  - authorization
  - rate limiting
- Better observability and access logging

Red flag

- One service reading another service's database

# Decentralized Data

What it means

- Each service owns:
  - its data
  - its schema
  - its storage technology

Why it matters

- Loose coupling
- Independent scaling and evolution

Security / audit perspective

- Stronger data isolation
- Different data classifications per service
- Reduced impact of a data breach

# What Is the 12-Factor App?

- Not a framework
- Not a tool
- Design methodology
- Cloud-native foundation

# Twelve-Factor App Principles

|    |                            |   |
|----|----------------------------|---|
| 1  | <b>Codebase</b>            | Application resources should have a version-managed, source-code repository                                 |
| 2  | <b>Dependencies</b>        | Should be set up in app manifest or configuration file and managed externally                               |
| 3  | <b>Config</b>              | Configurations of an app need to be stored independently as environment variables                           |
| 4  | <b>Backing Services</b>    | Services like databases, messaging systems, SMTP services, etc. should be architected as external resources |
| 5  | <b>Build, Release, Run</b> | There should be three independent steps of a deployment process   |
| 6  | <b>Stateless Processes</b> | Apps should have the provision to be served by multiple stateless, independent processes                    |
| 7  | <b>Port Binding</b>        | Apps directly bind to a port and responds to incoming requests  |
| 8  | <b>Concurrency</b>         | Apps need to be broken down into multiple modules for scaling   |
| 9  | <b>Disposability</b>       | Apps should maximize robustness with quick startup and easy shutdown  |
| 10 | <b>DEV/PROD Parity</b>     | Development, staging, and production should be kept as similar as possible                                  |
| 11 | <b>Logs</b>                | Treat logs as event streams   |
| 12 | <b>Admin Processes</b>     | Admin tasks should be run as one-off processes  |

# 12-Factor: Config

- No secrets in code
- Runtime injection
- Central secret management

# 12-Factor: Build, Release, Run

- Build = immutable artifact
- Release = artifact + config
- Run = execution only

# 12-Factor: Stateless & Logs

- Stateless services
- Logs as event streams
- Central logging

# 12-Factor vs Monolith (Audit)

- Better evidence
- Automation

# 12-Factor vs Monolith (Audit)

| Area     | Monolith | 12-Factor App        |
|----------|----------|----------------------|
| Config   | Files    | Env / Secret Manager |
| Logs     | Local    | Central              |
| Deploy   | Manual   | Automated            |
| Evidence | Weak     | Strong               |

# Microservice Patterns

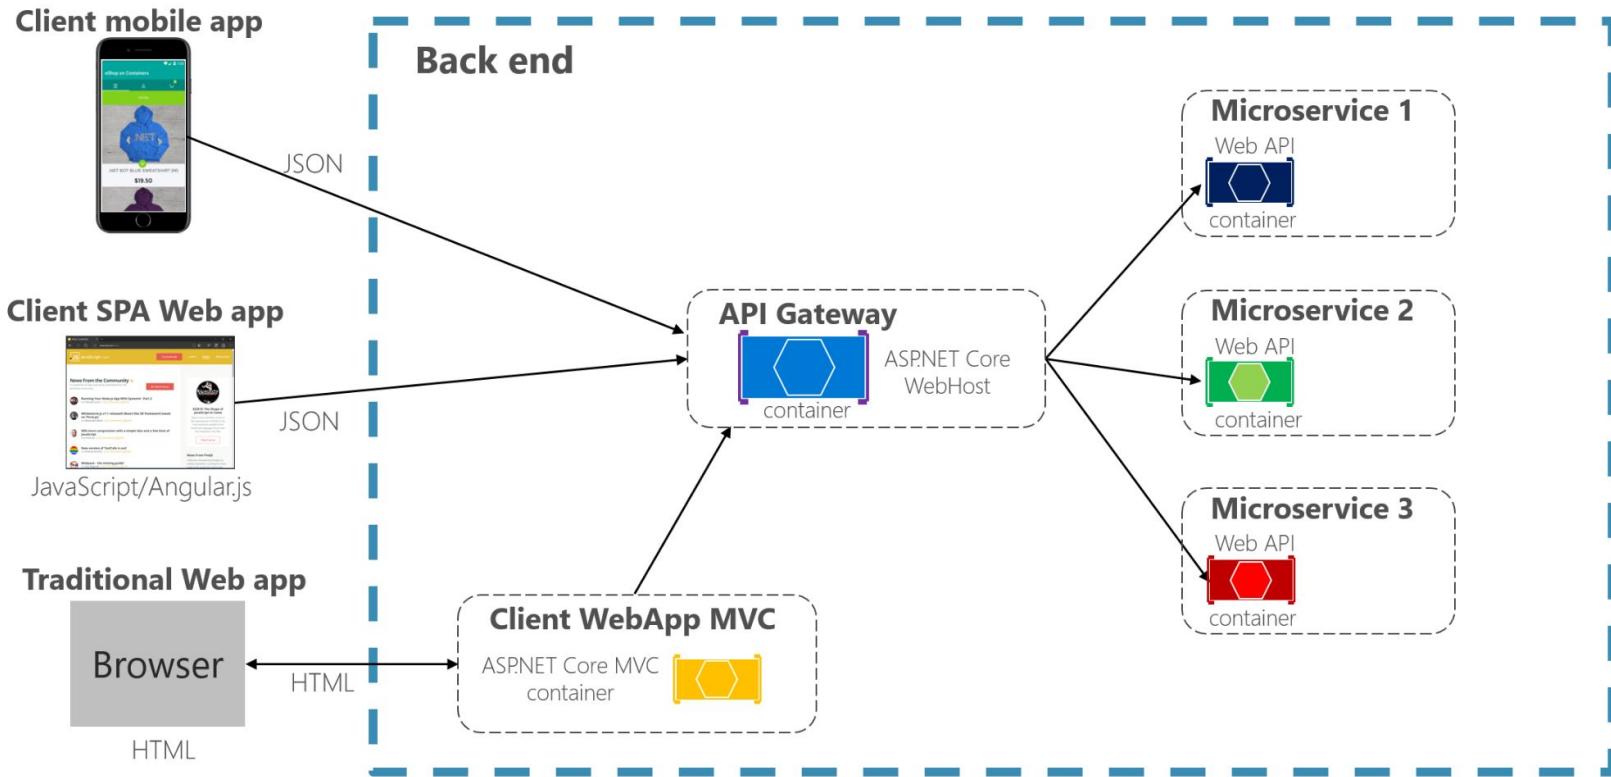
- Entry & Communication Patterns
  - API Gateway
  - Backend for Frontend
- Reliability & Resilience Patterns
  - Circuit Breaker
  - Retry with Backoff
  - Bulkhead
- Data Management Patterns
  - Database per Service
  - Saga
- Observability & Operations Patterns
  - Centralized Logging
  - Health Check
- Security Patterns
  - Zero Trust / mTLS (Service Mesh)

[More](#)

# API Gateway

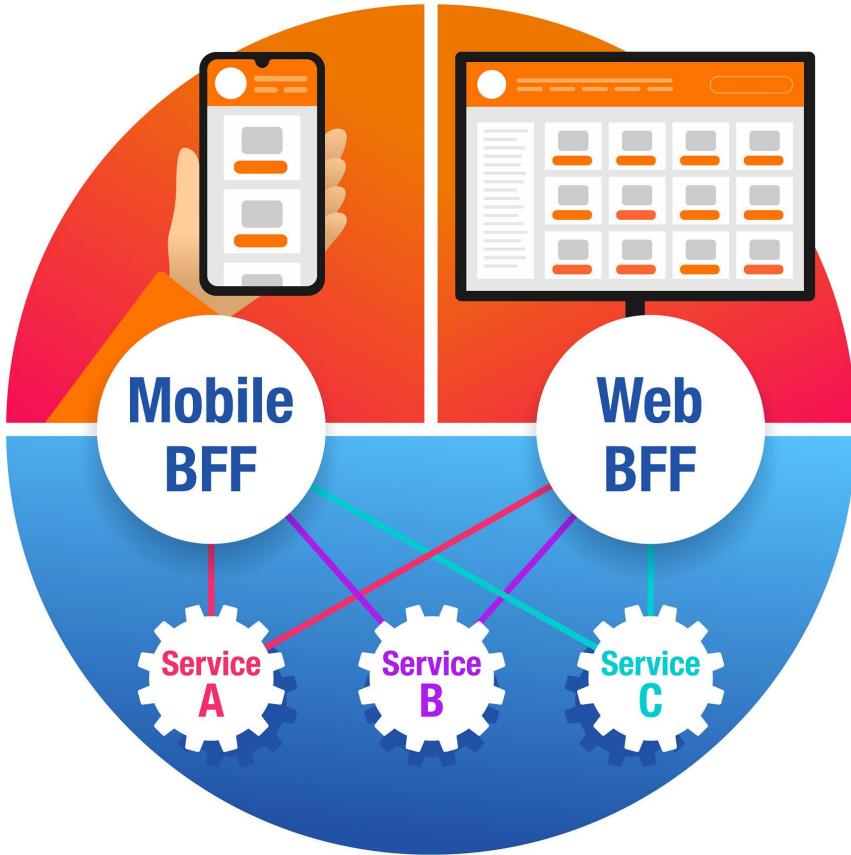
- What it is
  - Single entry point for external traffic to microservices.
- Pros
  - Centralized authentication & authorization
  - Rate limiting, request validation
  - Good audit control point
- Cons
  - Can become a bottleneck
  - High-value attack target
- When to use
  - Always for external access

# Using a single custom API Gateway service



# Backend for Frontend (BFF)

- What it is
  - Dedicated gateway per client type (web, mobile, partner API).
- Pros
  - Reduced attack surface per client
  - Clear separation of access rules
- Cons
  - More services to maintain
- When to use
  - Multiple clients with different needs



# Circuit Breaker

## What it is

- Stops requests to failing services.

## Pros

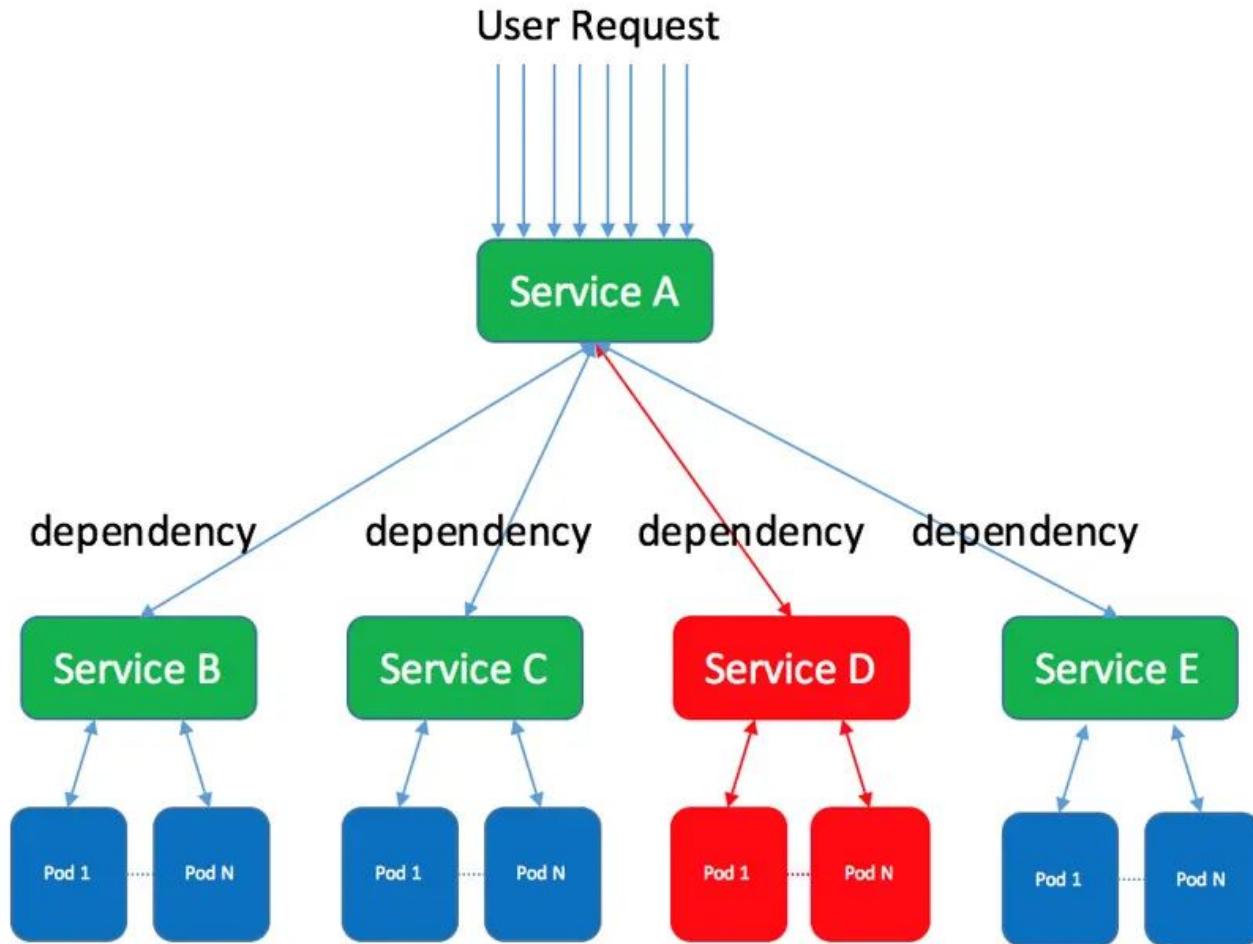
- Prevents cascading failures
- Improves system stability

## Cons

- Requires careful tuning

## When to use

- Systems with many service dependencies



# Retry with Backoff

## What it is

- Retries failed requests with increasing delays.

## Pros

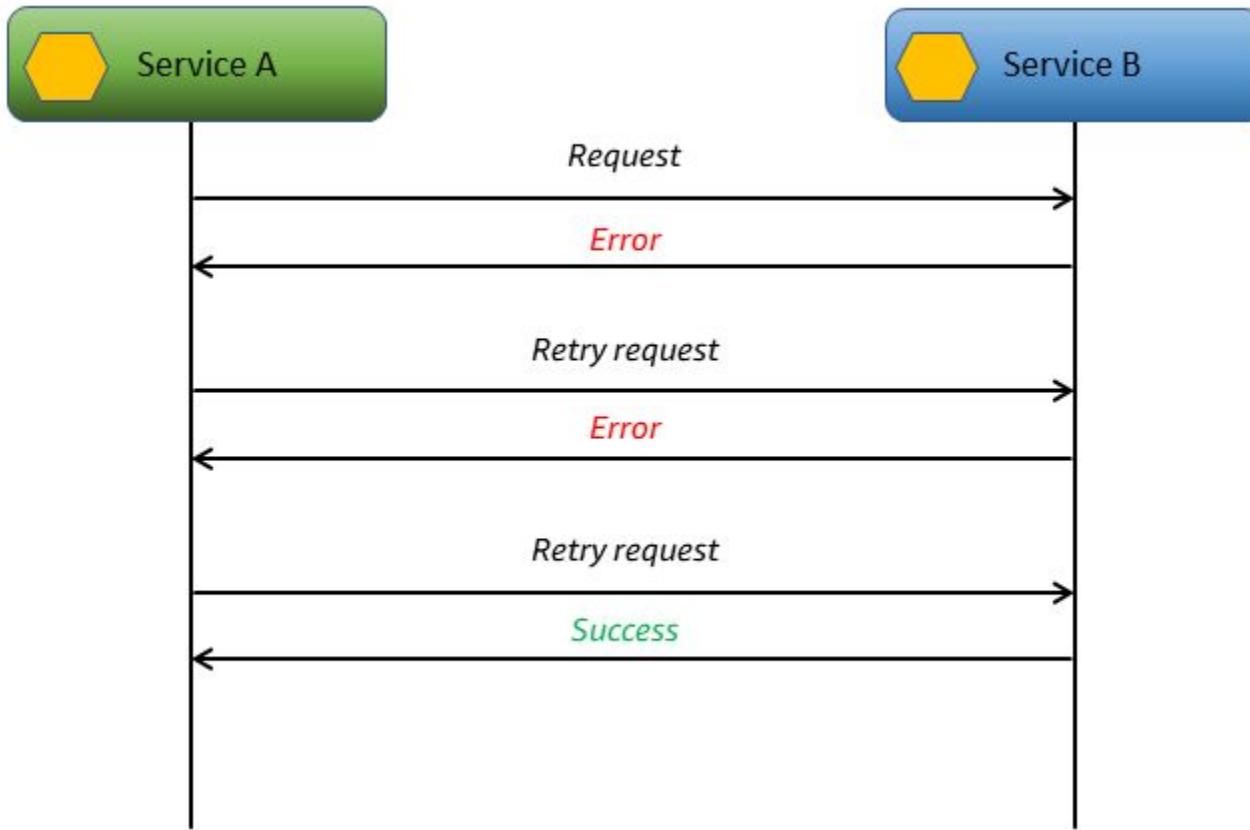
- Handles transient failures

## Cons

- Can amplify load if misused

## When to use

- Idempotent operations only



# Bulkhead

## What it is

- Resource isolation per service or dependency.

## Pros

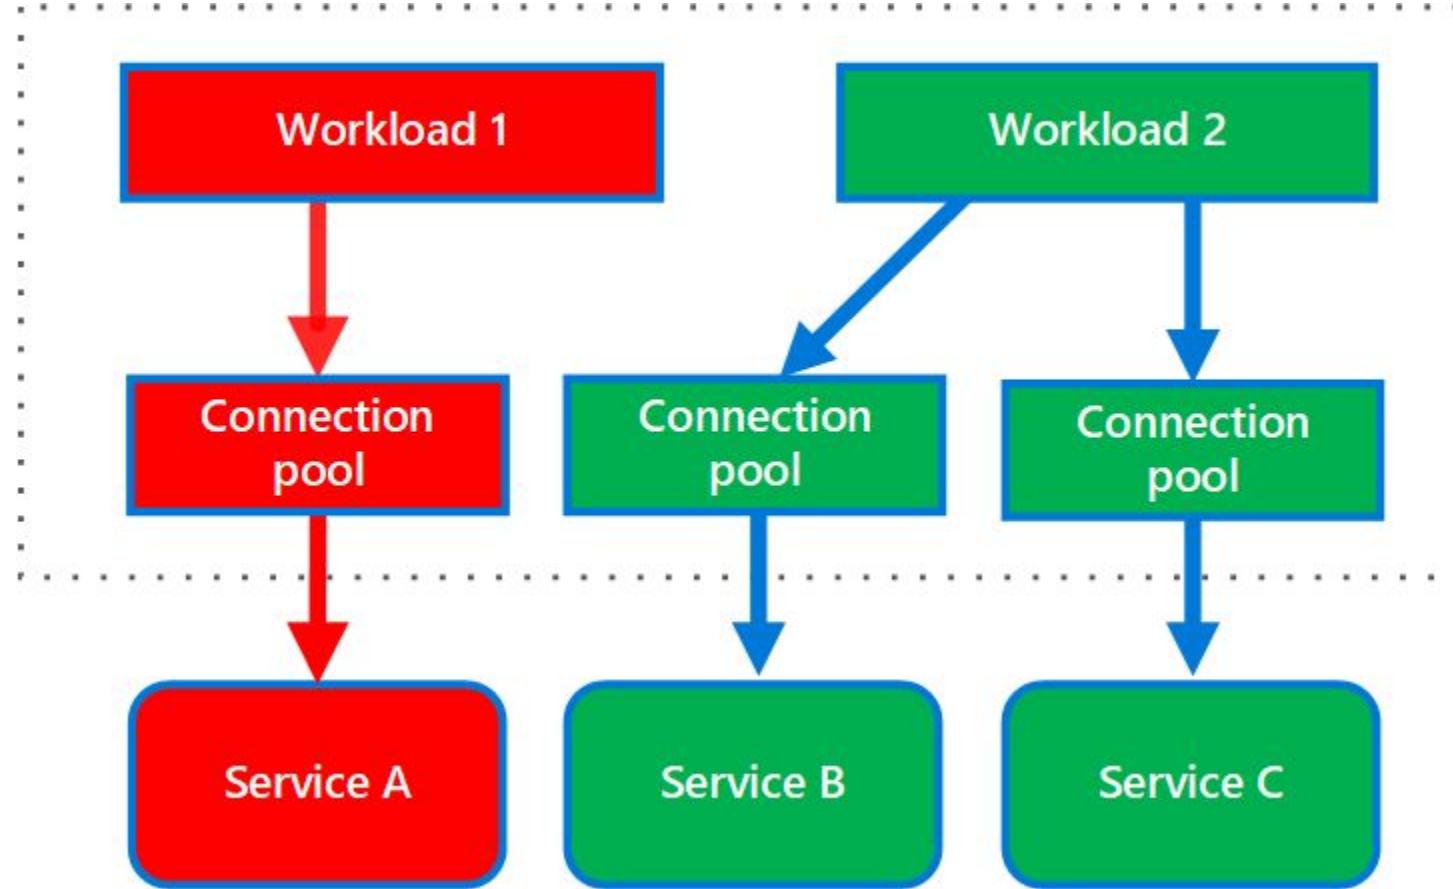
- Limits blast radius
- Improves fault isolation

## Cons

- Resource fragmentation

## When to use

- Critical or unstable dependencies



# Database per Service

## What it is

- Each service owns its database.

## Pros

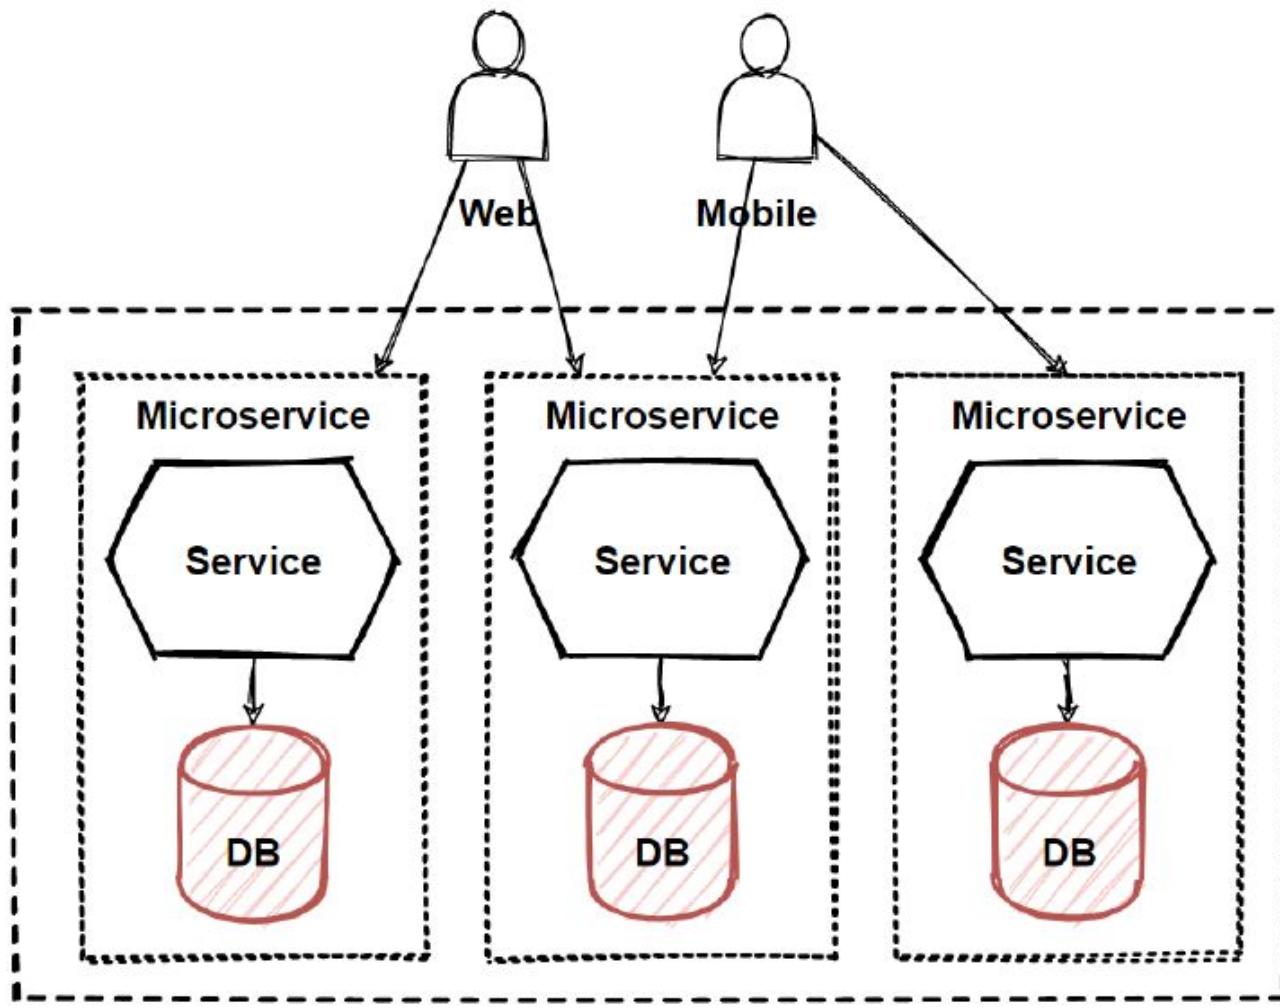
- Strong data isolation
- Independent scaling

## Cons

- Harder cross-service queries

## When to use

- Core microservice architectures



# Saga

## What it is

- Manages distributed transactions using events or orchestration.

## Pros

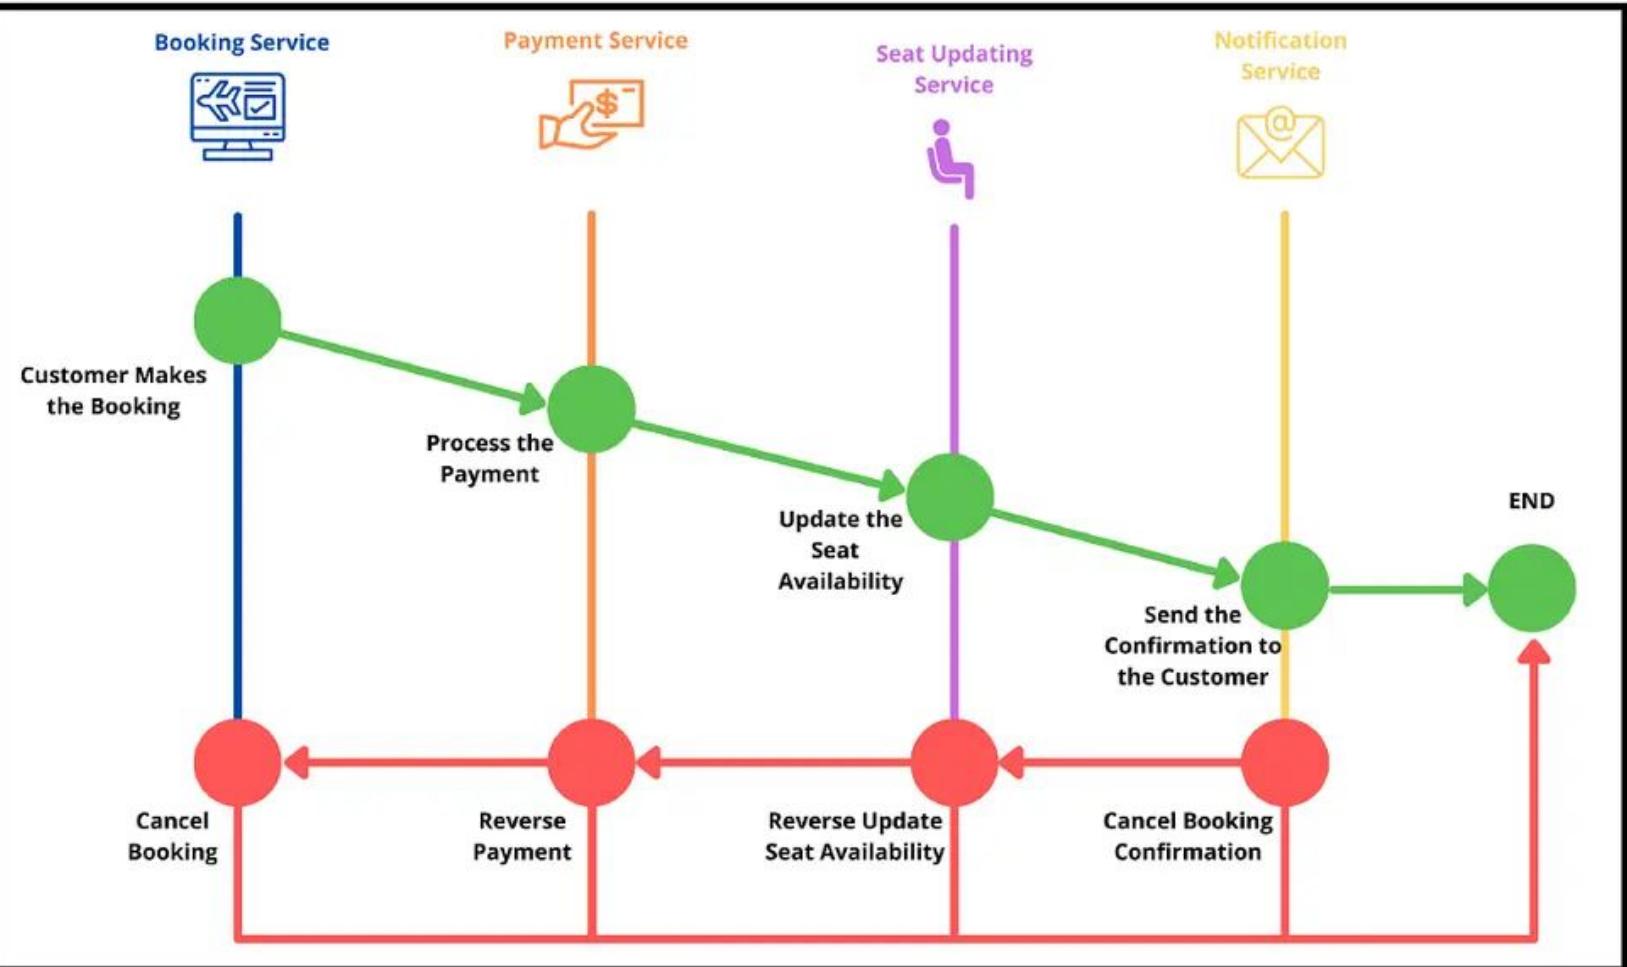
- No global locks
- Clear business workflows

## Cons

- Complex error handling

## When to use

- Multi-step business transactions



# Centralized Logging

## What it is

- Logs streamed to a central system.

## Pros

- Forensics & compliance
- Strong audit evidence

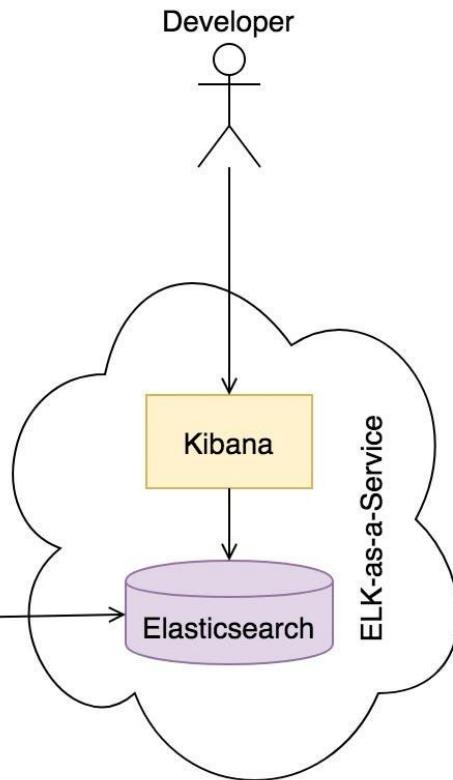
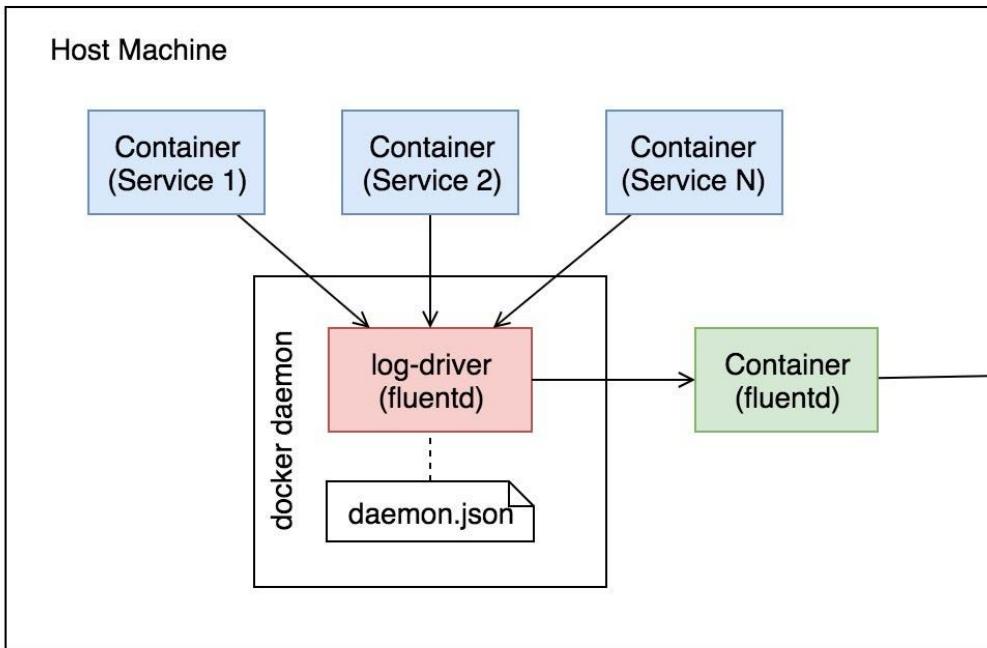
## Cons

- Storage & retention cost

## When to use

- Always (mandatory in regulated environments)

## Centralized Logging in Containerized Microservices



# Health Check (Liveness / Readiness)

## What it is

- Endpoints exposing service health.

## Pros

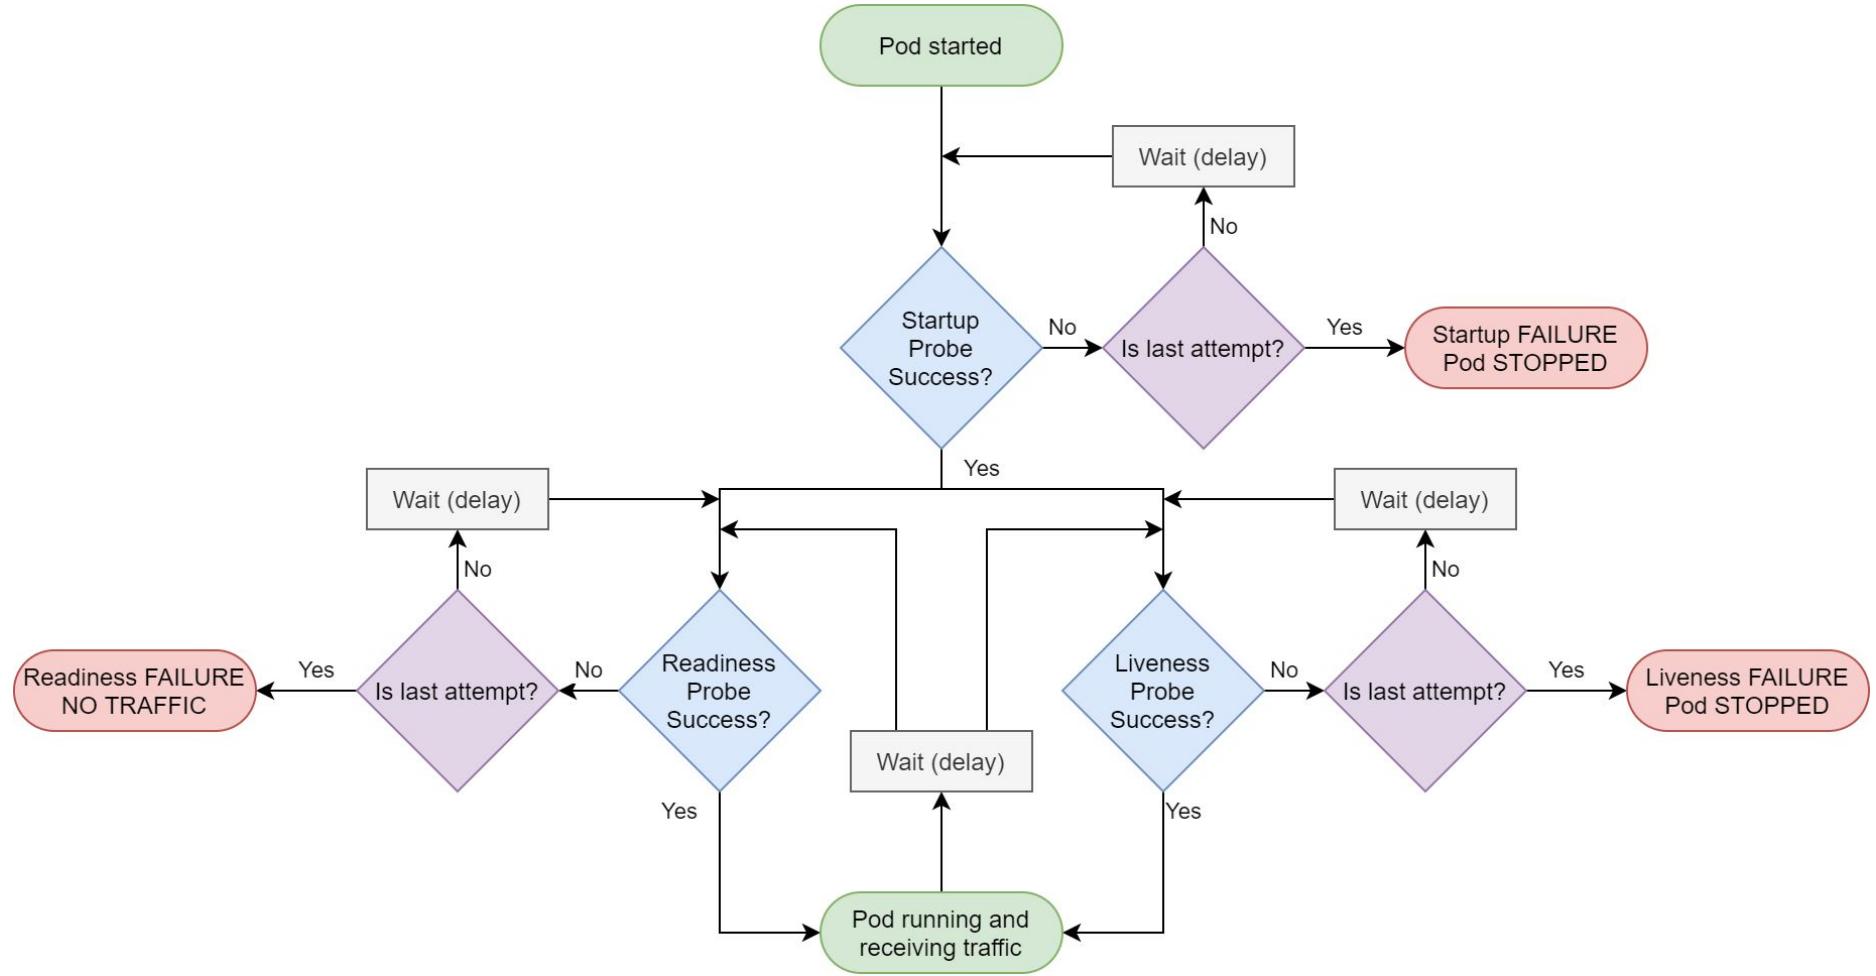
- Enables auto-healing
- Safe rollouts

## Cons

- Poor checks give false confidence

## When to use

- Every production service



# Service Mesh (mTLS / Zero Trust)

## What it is

- Mutual TLS and identity-based communication between services.

## Pros

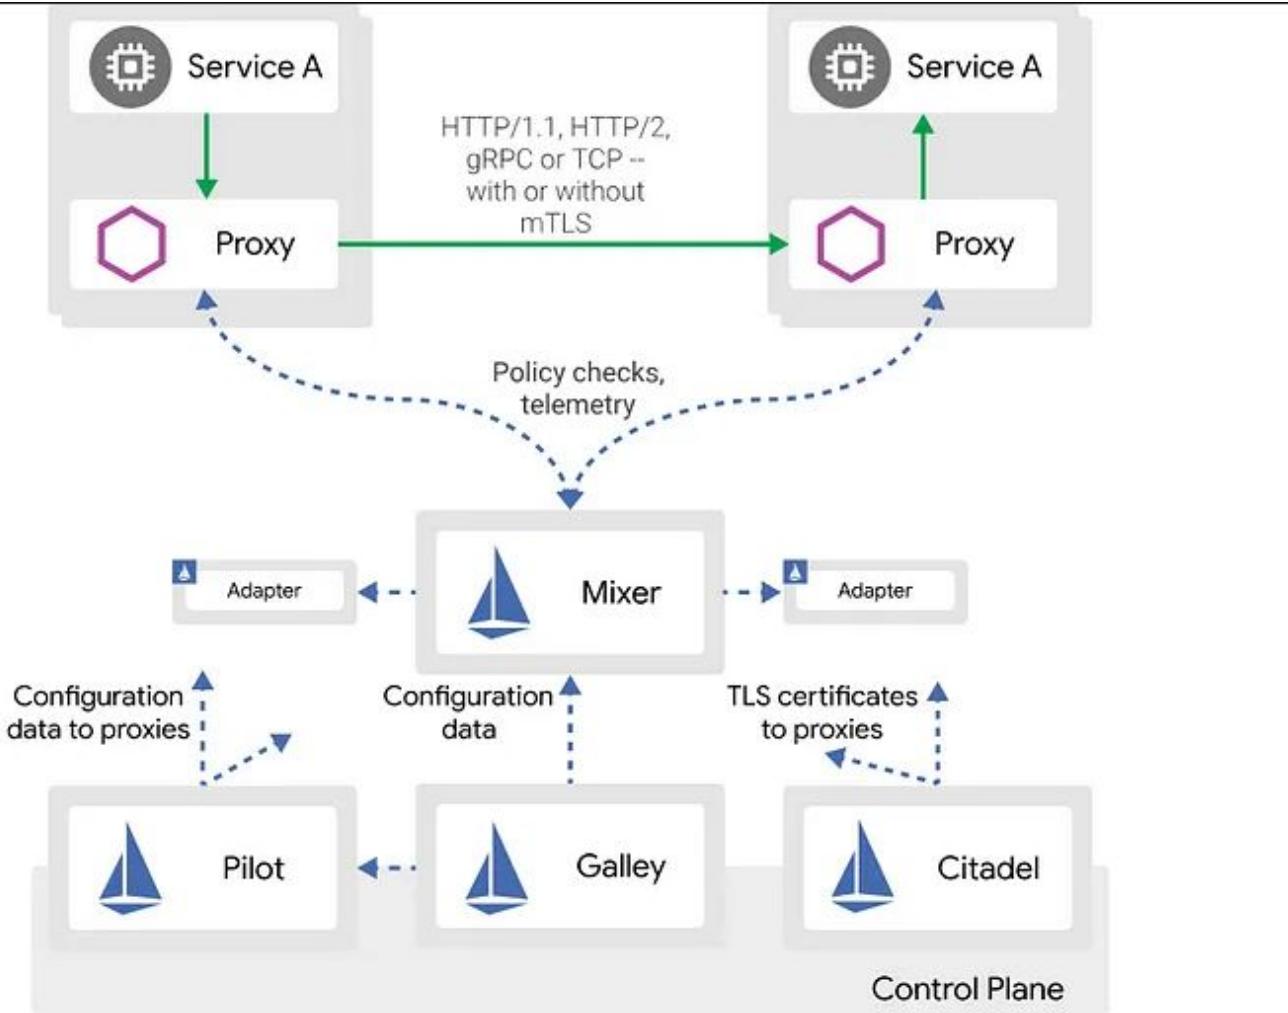
- Encrypted traffic
- Strong service identity
- Fine-grained authorization

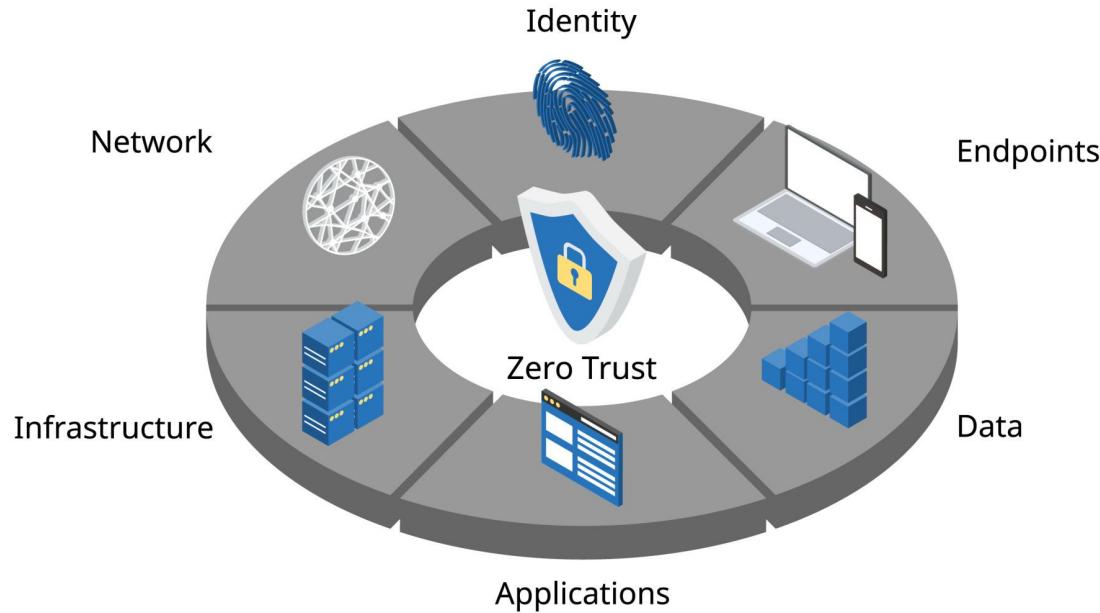
## Cons

- Operational complexity
- Performance overhead

## When to use

- Sensitive data
- Regulated environments



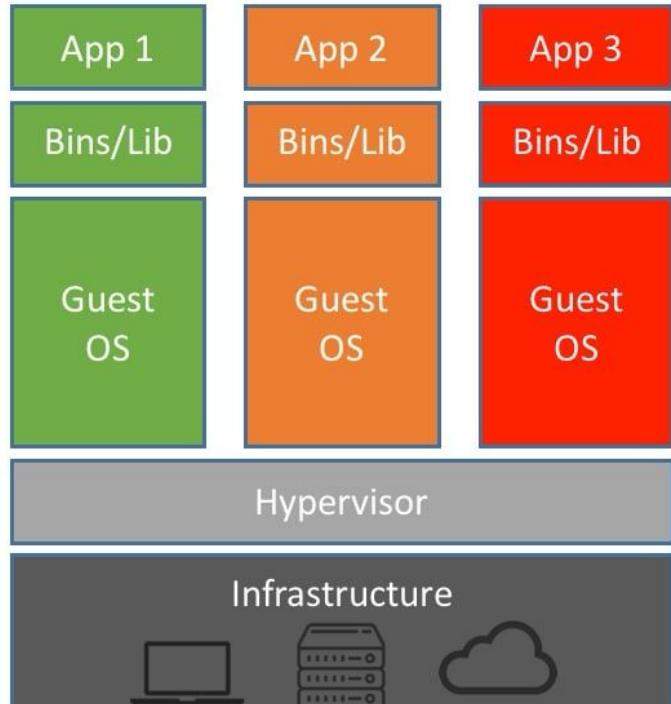


# **Containerization and Orchestration (Docker, Kubernetes)**

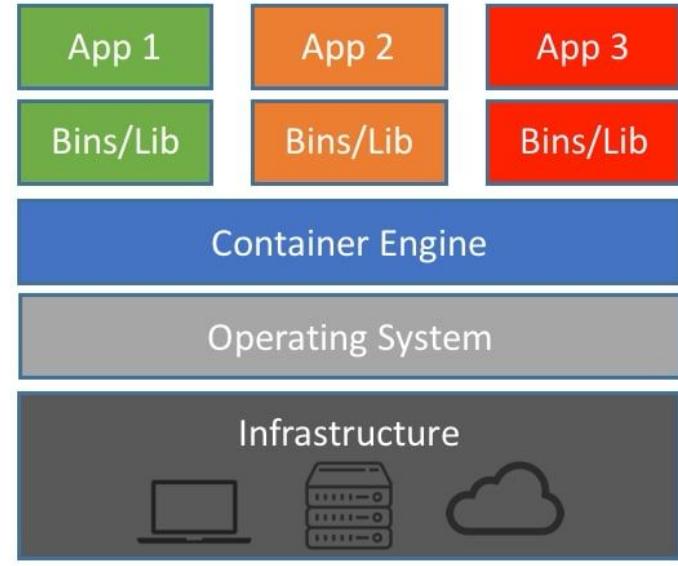
---

# What Are Containers?

- Containers are lightweight, isolated environments used to run applications
- They package:
  - application code
  - runtime
  - libraries and dependencies
- Containers share the host operating system kernel
- They start quickly and use fewer resources than virtual machines
- Containers are immutable at runtime (rebuilt, not modified)
- They are designed to be portable and consistent across environments



Machine Virtualization



Containers

# What Are Container Images?

- A container image is an immutable, versioned artifact used to run containers
- It contains:
  - application code
  - runtime
  - libraries and dependencies
- Images are typically stored in container registries
- An image represents a deployment unit and a supply-chain artifact
- Images are identified by:
  - tags (e.g. v1.2.3)
  - digests (SHA256 – preferred for security)

# What Are Container Images?

## Security relevance

- Images define what is actually running in production
- Vulnerabilities in images propagate to all running containers
- Image immutability enables rollback and auditability

# What Is a Dockerfile?

- A container image is an immutable, versioned artifact used to run containers
- It contains:
  - application code
  - runtime
  - libraries and dependencies
- Images are typically stored in container registries
- An image represents a deployment unit and a supply-chain artifact
- Images are identified by:
  - tags (e.g. v1.2.3)
  - digests (SHA256 – preferred for security)

# What Is a Dockerfile?

## Security relevance

- Images define what is actually running in production
- Vulnerabilities in images propagate to all running containers
- Image immutability enables rollback and auditability

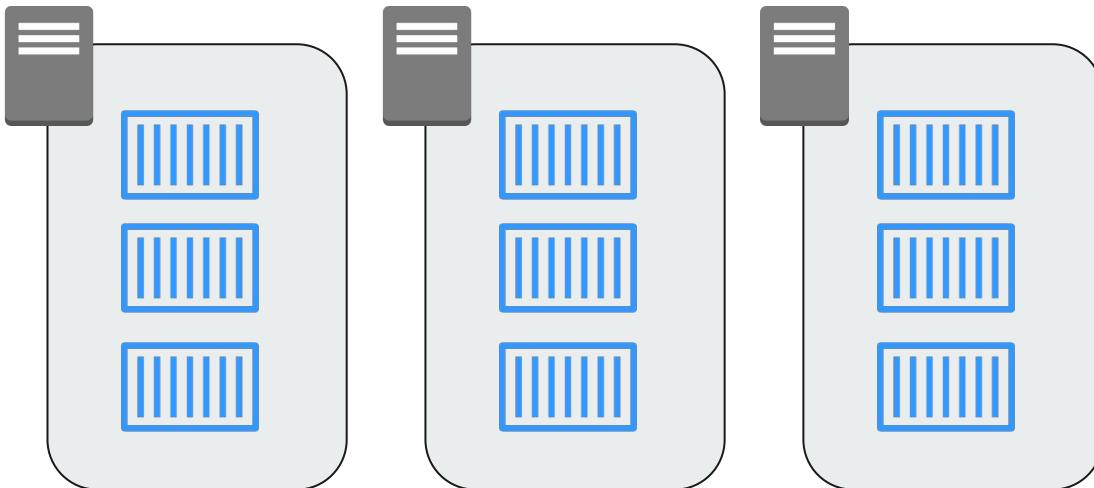
# Example Dockerfile

```
# Base image
FROM python:3.12-slim
# Working directory inside the container
WORKDIR /app
# Copy dependencies first (for caching)
COPY requirements.txt .
# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt
# Copy the application code
COPY . .
# Expose application port
EXPOSE 8080
# Define startup command
CMD ["python", "app.py"]
```

# What is Kubernetes?

“Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.” – CNCF

# The Problem Before Kubernetes

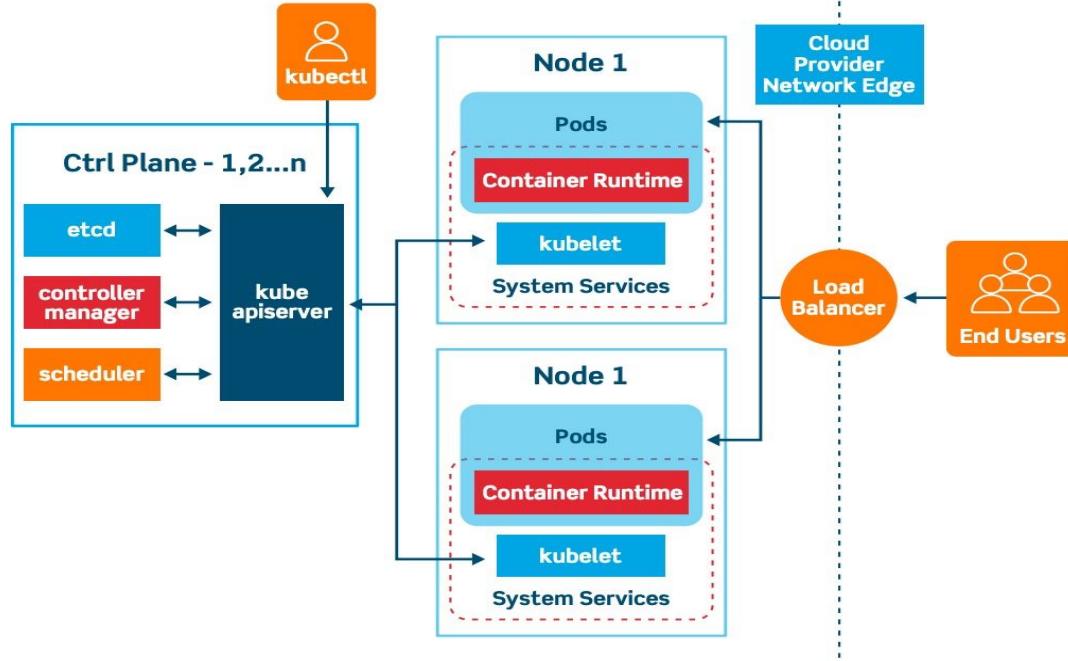


- Manual container management and updates
- Difficult scaling and recovery
- Updates can cause downtime
- Managing multiple containers was becoming difficult and time-consuming.
- No centralized control or observability

# Core Terminology

| Term       | Description                                    | Example                |
|------------|--|------------------------|
| Cluster    | A set of machines managed by Kubernetes        | Control + Worker nodes |
| Node       | A single machine (virtual or physical)         | Worker node            |
| Pod        | The smallest deployable unit                   | One or more containers |
| Service    | Stable access point to Pods                    | ClusterIP, NodePort    |
| Deployment | Controller that manages Pods and configuration | kubectl apply -f       |
| Namespace  | Logical resource grouping                      | dev, prod, staging     |

# Kubernetes Architecture



# API Server

- Acts as the central communication hub for the entire cluster.
- Exposes the Kubernetes REST API, which all components and clients (kubectl, controllers) use.
- Responsible for:
  - Validating and authenticating requests.
  - Persisting data in etcd.
  - Serving as the single source of truth for cluster state.

# etcd

- A distributed key–value database storing the full cluster configuration and state.
- Keeps both desired and current state of all objects.
- Critical component – regular backups are mandatory.
- Used by API Server for every read/write operation.

# Controller Manager

- Runs multiple internal controllers that constantly reconcile desired vs. actual state.
- Examples:
  - Node Controller – monitors node health.
  - ReplicaSet Controller – ensures correct number of Pod replicas.
  - Deployment Controller – manages rolling updates and rollbacks.
- If a Pod or Node disappears, the Controller Manager takes corrective action.

# Scheduler

- Decides which node should host a newly created Pod.
- Evaluates:
  - Resource availability (CPU, memory).
  - Node taints and tolerations.
  - Pod affinity and anti-affinity rules.
- Once a node is chosen, the Scheduler communicates the decision to the API Server.

# Cloud Controller Manager

- Integrates Kubernetes with external cloud providers (e.g., AWS, GCP, Azure).
- Manages:
  - Node lifecycle events (add/remove).
  - Load balancer provisioning.
  - Persistent volume integration.
- Makes Kubernetes cloud-aware while keeping Control Plane logic provider-agnostic.

# Kubelet

- Primary node agent that communicates directly with the API Server.
- Ensures all Pods assigned to the node are running and healthy.
- Core responsibilities:
  - Registers the node with the cluster.
  - Monitors Pod specs and reports status back to the Control Plane.
  - Restarts failed containers according to Pod definitions.
- Also handles liveness and readiness probes for health checking.

# kube-proxy

- Maintains the network connectivity rules inside the node.
- Implements Service load balancing and Pod-to-Pod routing.
- Works in one of two modes:
  - iptables mode – creates NAT rules for routing.
  - IPVS mode – high-performance connection-based load balancing.
  - nftables – modern replacement for iptables using the Linux kernel’s netfilter framework;
  - provides better performance, rule grouping, and scalability on new kernels.
- Ensures that traffic to a Service is evenly distributed among Pods.

# Container Runtime

- Responsible for running and managing containers defined in Pods.
- Implements the Container Runtime Interface (CRI).
- Common implementations:
  - containerd – lightweight, high-performance runtime (default on most distros).
  - CRI-O – optimized for Kubernetes, used by Red Hat and OpenShift.
  - (Docker Engine deprecated since K8s v1.24)
- Handles:
  - Image pulling, container lifecycle.
  - Resource isolation via cgroups and namespaces.

# Pod: The Smallest Deployable Unit

- The smallest deployable and schedulable unit in Kubernetes.
- Encapsulates one or more tightly coupled containers.
- Containers in the same Pod:
  - Share the same network namespace (same IP).
  - Can communicate via localhost.
  - Optionally share storage volumes.

# Init Containers

- Special containers that run before regular application containers.
- Execute one-time setup tasks, such as:
  - Waiting for a dependency (DB, API).
  - Initializing configuration or data.
  - Applying schema migrations.
- Each init container must complete successfully before the next one (or main containers) start.

# Sidecar

- A secondary container inside the same Pod as the main application.
- Adds supporting functionality without modifying the main app code.

# Adapter

- A secondary container inside the same Pod that transforms or normalizes data.
- Used when the main application outputs logs, metrics, or data in a non-standard format.
- The adapter translates it into a format that the rest of the platform understands.

# Ambasador

- A container in the same Pod that acts as a proxy or gateway for outbound connections.
- The main container connects to the ambassador over localhost.
- The ambassador handles communication with remote services — databases, APIs, or external networks.

# Deployment

- A Deployment is a Kubernetes object used to manage stateless applications
- It defines the desired state of Pods (how many, which image, how they are configured)
- Kubernetes automatically ensures the actual state matches the desired state
- A Deployment manages ReplicaSets, which in turn create and manage Pods
- It supports rolling updates and rollbacks
- It enables scaling (manual or automatic via HPA)
- It provides self-healing by recreating failed Pods
- Changes to a Deployment are versioned and auditable

# ConfigMap

- A Kubernetes object for storing non-sensitive configuration data.
- Key-value pairs such as:
  - feature flags,
  - log levels,
  - service URLs,
  - runtime tuning parameters.
- Lets you change behavior of an app without rebuilding the container image.

# Secret

- A Kubernetes object that stores small amounts of sensitive data - such as passwords, tokens, SSH keys, or certificates.
- Designed to separate confidential information from container images and configuration.
- Mounted or injected at runtime (never baked into images).

# Service: Stable Network Identity

- Provides a stable access endpoint for a dynamic set of Pods.
- Abstracts Pod IP changes using a virtual IP and DNS name.
- Implements load balancing across replicas.
- Common types:
  - ClusterIP – internal cluster access (default).
  - NodePort – exposes on node ports (30000–32767).
  - LoadBalancer – integrates with external cloud LB.

# Job

- A controller that runs Pods until they successfully complete.
- Guarantees that a defined number of Pods finish successfully.
- Retries automatically on failure (depending on configuration).
- Used for batch processing, migrations, and one-time tasks.

# CronJob

- Creates Jobs on a time-based schedule, like Linux cron.
- Each schedule execution spawns a new Job.
- Great for:
  - Backups
  - Reports
  - Database maintenance
  - Cleanup tasks

# Ingress

- Ingress is a Kubernetes object that defines HTTP/HTTPS routing rules
- It exposes services outside the cluster using:
  - hostnames
  - paths
- Ingress works at Layer 7 (HTTP)
- It supports:
  - TLS termination
  - host-based routing
  - path-based routing
- Ingress is only a configuration object, not a traffic handler

# What Is an Ingress Controller?

- An Ingress Controller is the component that:
  - watches Ingress resources
  - implements the routing rules
- It is a running application inside the cluster
- Common examples:
  - NGINX Ingress Controller
  - HAProxy
  - Cloud provider load balancer controllers
- Responsible for:
  - traffic forwarding
  - TLS handling
  - applying security features (rate limiting, headers)

# What Is an Ingress Controller?

Security / audit relevance

- Central enforcement point for:
  - TLS
  - authentication
  - request logging
- Misconfiguration can expose the entire cluster

# Gateway API

- Gateway API is the next-generation Kubernetes API for traffic management
- Designed to replace and extend Ingress
- More expressive and more structured
- Built around clear roles and responsibilities
- Key design goals
  - Standardized
  - Extensible
  - Role-oriented (platform vs application teams)

# Probes, Autoscaling, Rollouts

- Liveness vs readiness
- HPA, VPA
- Canary / Blue-Green

# What are probes?

- Probes are health checks executed by Kubernetes against containers
- They determine whether a container is healthy and ready to receive traffic

# Liveness Probe

## What it does

- Checks if the application is still running correctly
- If it fails repeatedly → container is restarted

## Why it matters

- Detects deadlocks or stuck processes
- Enables self-healing

## Security / audit perspective

- Prevents long-running faulty or compromised processes
- Provides evidence of automated remediation

## Common mistake

- Using liveness checks for dependency health (DB, API)

# Readiness Probe

## What it does

- Checks if the application is ready to accept traffic
- If it fails → Pod is removed from Service endpoints
- Container is not restarted

## Why it matters

- Zero-downtime deployments
- Safe startup and graceful degradation

## Security / audit perspective

- Reduces exposure to half-initialized or misconfigured services
- Limits blast radius during incidents

# What is autoscaling?

- Autoscaling dynamically adjusts the number of Pods
- Based on:
  - CPU
  - memory
  - custom or external metrics

# Horizontal Pod Autoscaler (HPA)

## What it does

- Scales Pods up or down automatically
- Reacts to load in near real time

## Why it matters

- Handles traffic spikes without manual intervention
- Improves availability and resilience

## Security implications

- Reduces DoS impact by absorbing traffic
- Prevents single-instance overload
- Requires quotas and limits to avoid cost abuse

## Audit evidence

- HPA configuration
- Metrics history
- Scaling events (who, when, why)

# Vertical Pod Autoscaler (VPA)

- Vertical Pod Autoscaler (VPA) automatically adjusts:
  - CPU requests
  - memory requests
  - It optimizes resource allocation per Pod
  - VPA focuses on right-sizing, not scaling out
- Works best for:
  - stable workloads
  - batch jobs
  - services with predictable traffic

# What is a rollout?

- A rollout is the process of deploying a new application version
- Managed automatically by Kubernetes Deployments
- By default Kubernetes only support RollingUpdate and Replace other like Canary or A/B need to be handle externally

# Rolling Update (default)

## What it does

- Gradually replaces old Pods with new ones
- Keeps the application available during updates

## Why it matters

- No downtime
- Reduced deployment risk
- Security / audit perspective

## Each rollout:

- is versioned
- is traceable
- can be rolled back
- Strong change-management evidence

# What “Security” Mean in Kubernetes?

Kubernetes security covers multiple layers, not just user access.

It involves securing:

- The API Server — who can talk to the control plane.
- Workloads — what containers and Pods can do.
- The Network — which Pods can talk to each other.
- The Supply Chain — ensuring trusted images and manifests.

# Kubernetes Security Responsibilities

Developers:

- Avoid running containers as root
- Limit container privileges
- Store secrets safely

Cluster Operators:

- Manage authentication & RBAC
- Apply PodSecurity standards
- Enforce admission policies

Infrastructure / Platform Team:

- Secure API server and etcd
- Patch the control plane
- Configure audit logs and encryption

# ServiceAccount

A ServiceAccount (SA) is an identity used by processes running inside the cluster to authenticate to the Kubernetes API.

- Every Pod runs under a ServiceAccount context.
- ServiceAccounts are meant for machines, Pods, or controllers — not human users.
- They're essential for fine-grained access control via RBAC.

# Continuous Integration & Continuous Delivery/Deployment

---

---

# What is CI/CD

# Continuous Integration (CI)

## What it means

- Developers frequently merge code into a shared repository
- Each change automatically triggers:
  - build
  - tests
  - static analysis

## Why it matters

- Early detection of bugs and misconfigurations
- Prevents broken or insecure code from progressing further

# Continuous Integration (CI)

## Security / audit perspective

- CI enforces shift-left security
- Security checks become:
- repeatable
- mandatory
- logged
- Eliminates “it worked on my machine”

## Audit evidence

- Pipeline execution logs
- Test results
- Static analysis reports
- Commit -> pipeline traceability

# Continuous Delivery/Deployment (CD)

## What it means

- Continuous Delivery: code is always deployable, release is manual
- Continuous Deployment: code is deployed automatically

## Why it matters

- Fast and controlled releases
- Reduced human error during deployments

# Continuous Delivery/Deployment (CD)

## Security / audit perspective

- Deployment logic is code, not a human procedure
- Clear separation:
  - who can change code
  - who can approve release
- Rollbacks are automated and logged

## Audit evidence

- Deployment history
- Approval records
- Environment-specific policies
- Git tags / release metadata

---

# CI/CD Tools Landscape

# Source Control

- Git repositories are the single source of truth
- Every change is:
  - versioned
  - attributed to an identity
  - reviewable

## Audit relevance

- Strong non-repudiation
- Full change history

# CI Systems (Build & Test)

Examples:

- GitHub Actions
- GitLab CI
- Jenkins

What they do

- Build artifacts
- Run tests
- Enforce quality and security checks

Audit relevance

- Evidence of:
  - automated validation
  - policy enforcement
  - failed vs successful runs

# CD / GitOps Tools (Deploy)

Examples:

- ArgoCD
- Flux

What they do

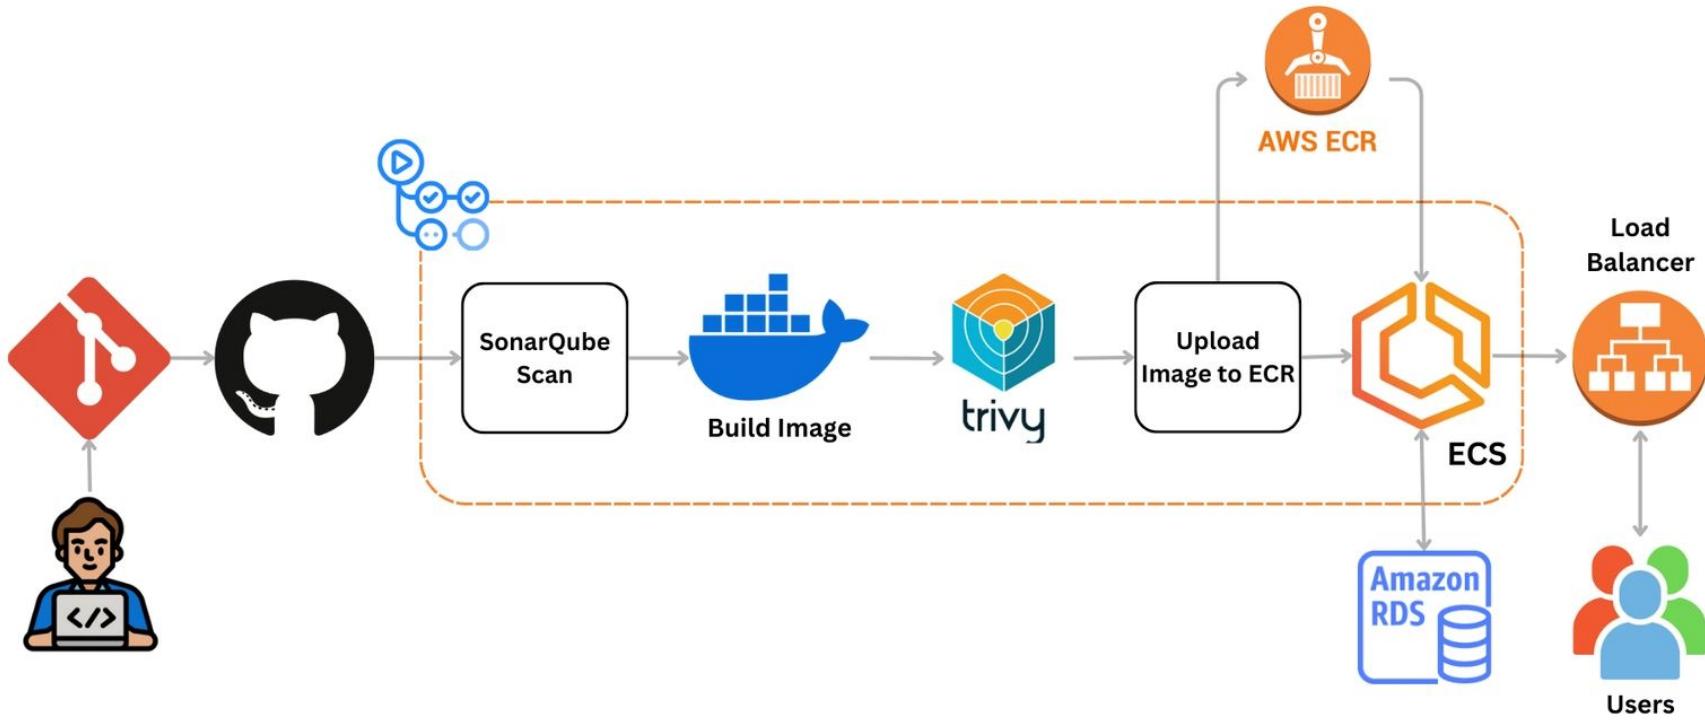
- Reconcile desired state from Git
- Continuously enforce configuration

Audit relevance

- No “manual deploys”
- Git becomes deployment evidence

# Common audit red flags

- Pipelines editable without review
- Secrets stored in pipeline YAML
- Manual production deploys via CLI
- No artifact immutability



**Dev** stores code in local SCM (**Git**), pushes to the central repository (**Github**)

**Github actions** is configured to run the pipeline after a **commit** is detected. It first runs a static code scan with **Sonar Cloud**, then builds the **Docker image**, uses **Trivy** to scan for vulnerabilities and **uploads** the image to **ECR**

**ECR** picks the **uploaded** image, the pipeline passes the recently **built version** to ECR which then updates its **environment**.

Users access the **application** through a **load balancer**

---

# Analyze Pieline

# Infrastructure as Code



# What Is Infrastructure as Code?

- Infrastructure is defined using code, not manual configuration
- The desired state of infrastructure is:
  - declarative
  - version-controlled
  - reproducible
- Changes are applied through automated workflows

# Why IaC Matters

## Operational perspective

- Consistent environments (dev = prod)
- Fast provisioning and teardown
- Reduced configuration drift

## Security & audit perspective

- Infrastructure changes follow the same controls as application code
- Enables:
  - peer reviews
  - approvals
  - rollback
- Eliminates “click-ops” (manual console changes)

# Core IaC Tools

- Terraform – cloud resources
- Helm – application configuration
- Kustomize – environment overlays

---

# Auditability of IaC

# What Makes IaC Auditable?

- Version Control
- Plan vs Apply
- Drift Detection

# Version Control

Every change is:

- tracked
- attributed to an identity
- timestamped

Audit evidence

- Git commit history
- Pull requests
- Code reviews

# Plan vs Apply

- terraform plan shows what will change
- terraform apply executes approved changes

## Audit relevance

- Separation between:
  - intent
  - execution
- Strong change-management control

# Drift Detection

Detects:

- manual changes
- configuration violations
- Ensures runtime state matches declared state

Audit relevance

- Proof that manual changes are not tolerated
- Evidence of continuous compliance

# IaC Security & Compliance Tooling

| Tool                       | Purpose                     | Typical Use                    |
|----------------------------|-----------------------------|--------------------------------|
| Checkov                    | Policy-as-code scanning     | Prevent misconfigurations      |
| tfsec                      | Terraform security scanning | Fast local checks              |
| Terraform Cloud/Enterprise | Policy enforcement          | Governance & approvals         |
| OPA / Conftest             | Custom compliance rules     | Organization-specific policies |
| Infracost                  | Cost impact analysis        | Financial risk control         |

# Logging, Monitoring and Observability

---

---

# Observability

# Logs, Metrics, Traces

## Logs

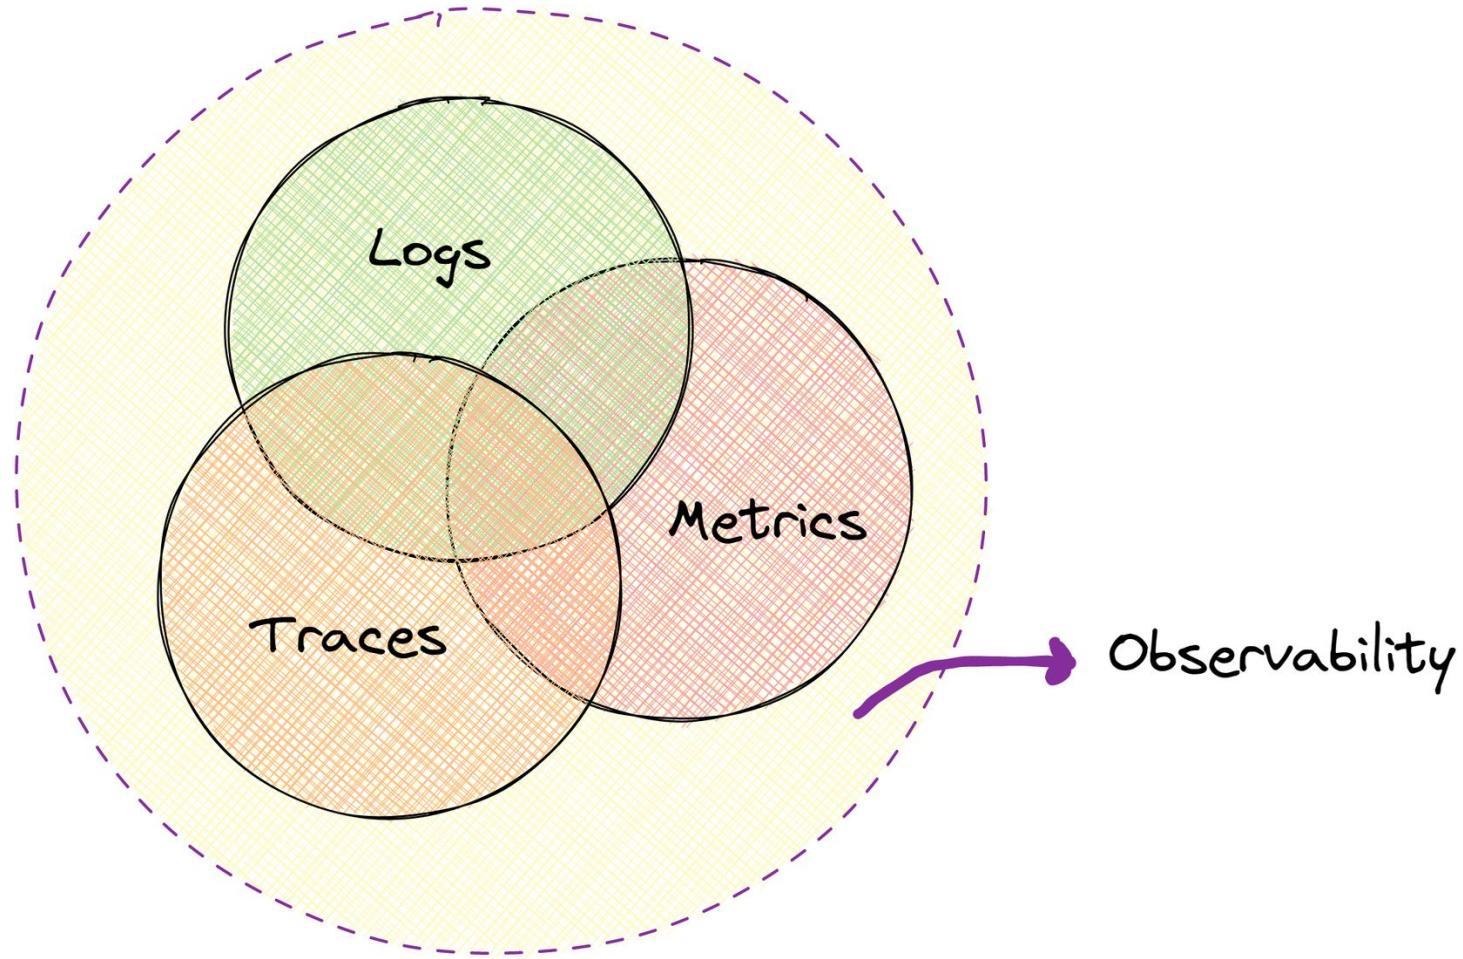
- Discrete, timestamped records of events
- Human-readable or structured
- Best for forensics and investigations

## Metrics

- Numeric measurements over time
- Aggregated and efficient
- Best for health, trends, and alerting

## Traces

- End-to-end view of a request
- Show how a request flows across services
- Best for root-cause analysis



# Why Observability Matters

## Operational view

- Faster incident detection
- Faster recovery
- Reduced mean time to resolution (MTTR)

## Security & audit view

- Visibility into:
  - system behavior
  - access patterns
  - failures and anomalies
- Enables evidence based conclusions, not assumptions

---

# Log Retention & Access

# Log Retention

## What it means

- How long logs are stored
- Where they are stored
- In which form (raw vs aggregated)

## Why it matters

- Compliance requirements (e.g. months vs years)
- Ability to investigate incidents after the fact

# Log Retention

## Audit considerations

- Retention must be:
- defined
- documented
- enforced automatically

# Log Access Control

## What it means

- Who can:
  - read logs
  - export logs
  - delete logs

## Security risks

- Logs often contain:
  - IP addresses
  - user identifiers
  - request metadata
- Logs can be sensitive data

# Log Access Control

## Audit evidence

- RBAC policies for logging systems
- Access logs for log systems themselves
- Immutable or write-once storage (where required)

# Common Audit Red Flags

- Logs stored locally on nodes
- No defined retention period
- Shared admin access to logs
- Logs deletable without trace

---

# Tracing and Metrics: Audit Value

# Why Tracing Has Audit Value

Shows how a request moved through the system

- Identifies:
  - which services were involved
  - where failures occurred
  - how long each step took

Audit use cases

- Incident reconstruction
- Data-flow verification
- Proof of segregation between services

# Why Metrics Have Audit Value

Provide objective, quantitative evidence

- Answer questions like:
  - Was the system overloaded?
  - Did autoscaling work as expected?
  - Were SLAs violated?

Audit use cases

- Capacity management validation
- DoS and abuse analysis
- Operational risk assessment

---

# Prometheus, Grafana, ELK / EFK Stacks

# Prometheus

## What it is

- Metrics collection and storage system
- Pull-based model
- Native Kubernetes integration

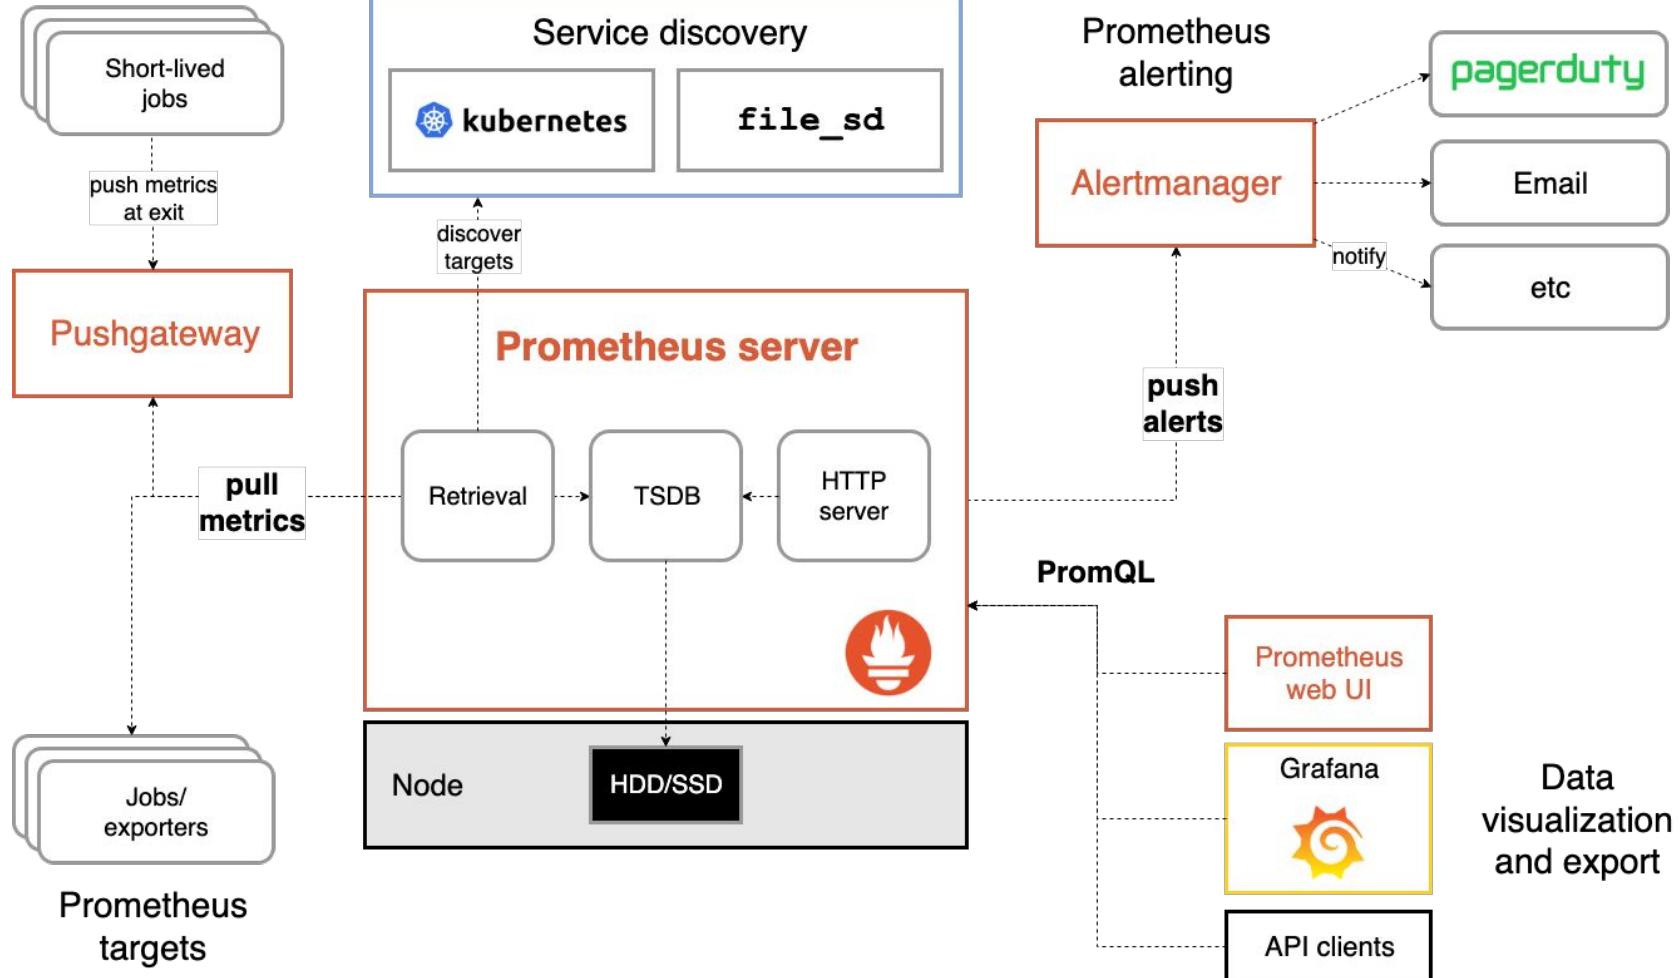
## Typical use

- CPU, memory, request rates
- Autoscaling input (HPA)
- Alerting

# Prometheus

## Audit value

- Historical performance data
- Evidence of capacity controls
- Proof of SLA / SLO monitoring



# Grafana

## What it is

- Visualization and dashboarding tool
- Works with Prometheus, Elasticsearch, and more
- Typical use

## Dashboards for:

- system health
- application performance
- security signals

# Grafana

## Audit value

- Read-only dashboards for auditors
- Visual evidence of controls in action
- Historical views for specific incidents



# ELK / EFK Stack

## What it is

- Elasticsearch — log storage and search
- L/F — Logstash or Fluentd — log ingestion
- Kibana — visualization and querying

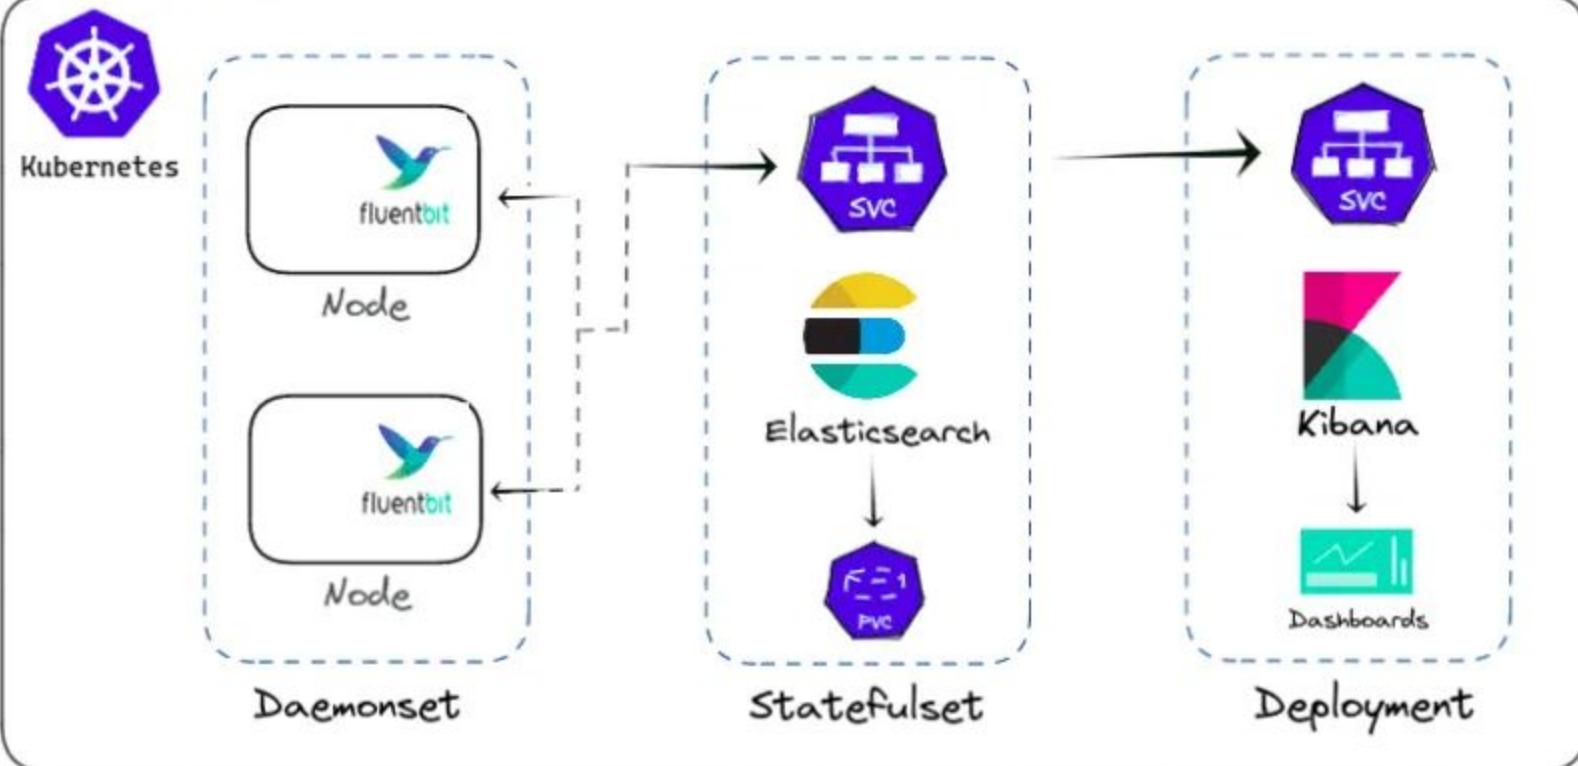
## Typical use

- Centralized logging
- Full-text search
- Correlation across systems

# ELK / EFK Stack

## Audit value

- Forensic investigations
- Long-term log retention
- Advanced filtering and correlation



# Security in Cloud Native Environments

---

---

# Kubernetes RBAC

# What is Kubernetes RBAC?

- RBAC (Role-Based Access Control) defines who can do what in a Kubernetes cluster
- Permissions are granted via:
  - Roles / ClusterRoles (what actions are allowed)
  - RoleBindings / ClusterRoleBindings (who gets those permissions)

# Core RBAC Objects

- Role — permissions within a namespace
- ClusterRole — cluster-wide permissions
- RoleBinding — binds a Role to a user/service account
- ClusterRoleBinding — binds a ClusterRole globally

# Why RBAC Matters

## Security perspective

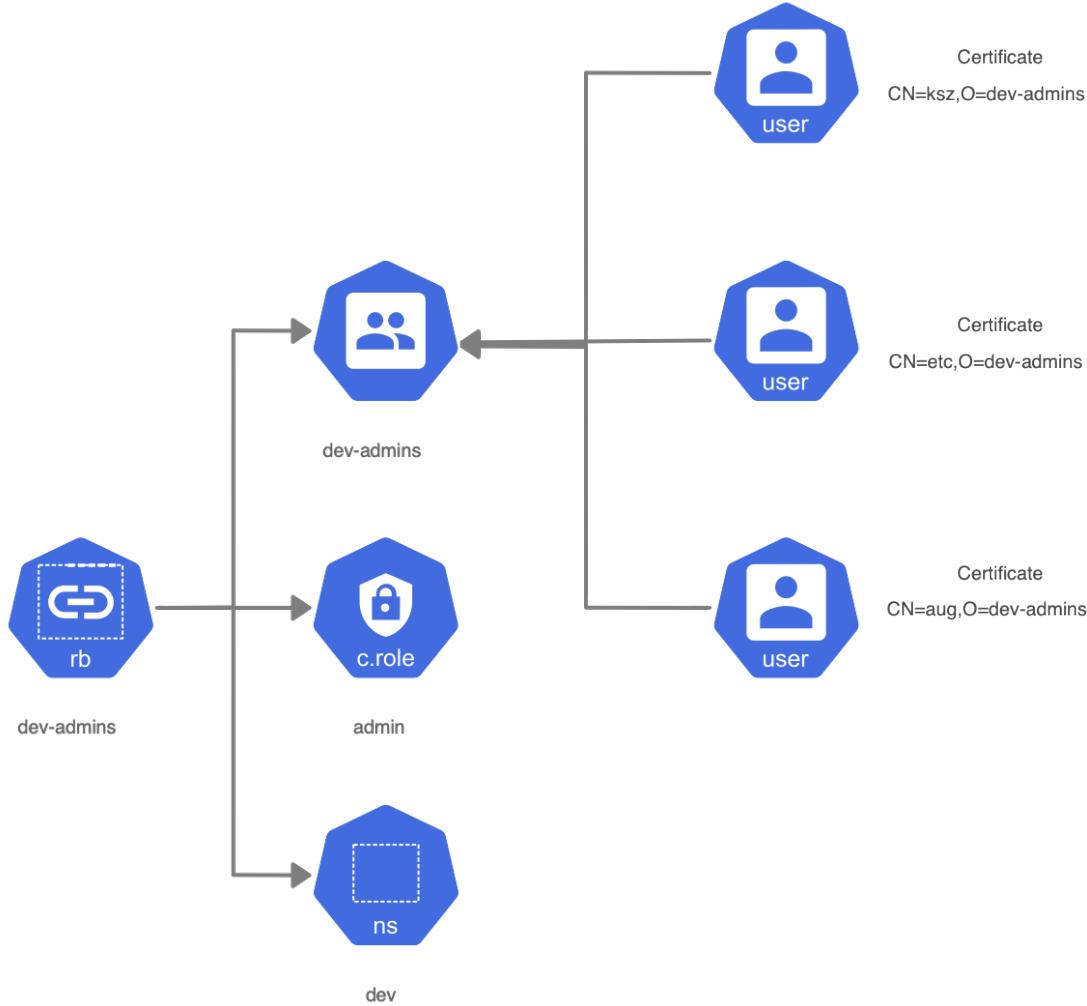
- Enforces least privilege
- Limits blast radius of compromised identities
- Separates human and workload access

## Audit perspective

- Clear, declarative access rules
- Easy to answer:
  - Who can deploy?
  - Who can read secrets?
- RBAC configs = direct audit evidence

# Common RBAC Anti-Patterns

- cluster-admin used by applications
- Shared service accounts
- RBAC changes outside Git
- No periodic access reviews



---

# Secret Management

# Kubernetes Secrets (Native)

- Store sensitive data (tokens, passwords, keys)
- Mounted as:
  - environment variables
  - files
- Base64-encoded, not encrypted by default

# External Secret Management

- Centralized secret storage
- Dynamic secrets (short-lived credentials)
- Automatic rotation and revocation
- Strong audit logging

# Audit Questions to Ask

- Where are secrets stored?
- Who can access them?
- How often are they rotated?
- Can access be revoked immediately?

# External Secret Management Integration

Google Cloud Secret Manager



sveltos-secret



---

# Image scanning, digital signatures, SBOMs

# Image Scanning

Scans container images for:

- known vulnerabilities (CVEs)
- outdated libraries
- Typically integrated into CI pipelines

Audit value

- Proof that images are checked before deployment
- Risk-based acceptance of vulnerabilities

# Digital Signatures (Image Signing)

- Cryptographically sign container images
- Verify image integrity and origin at deploy time

## Why it matters

- Prevents tampering
- Ensures only trusted images run in production

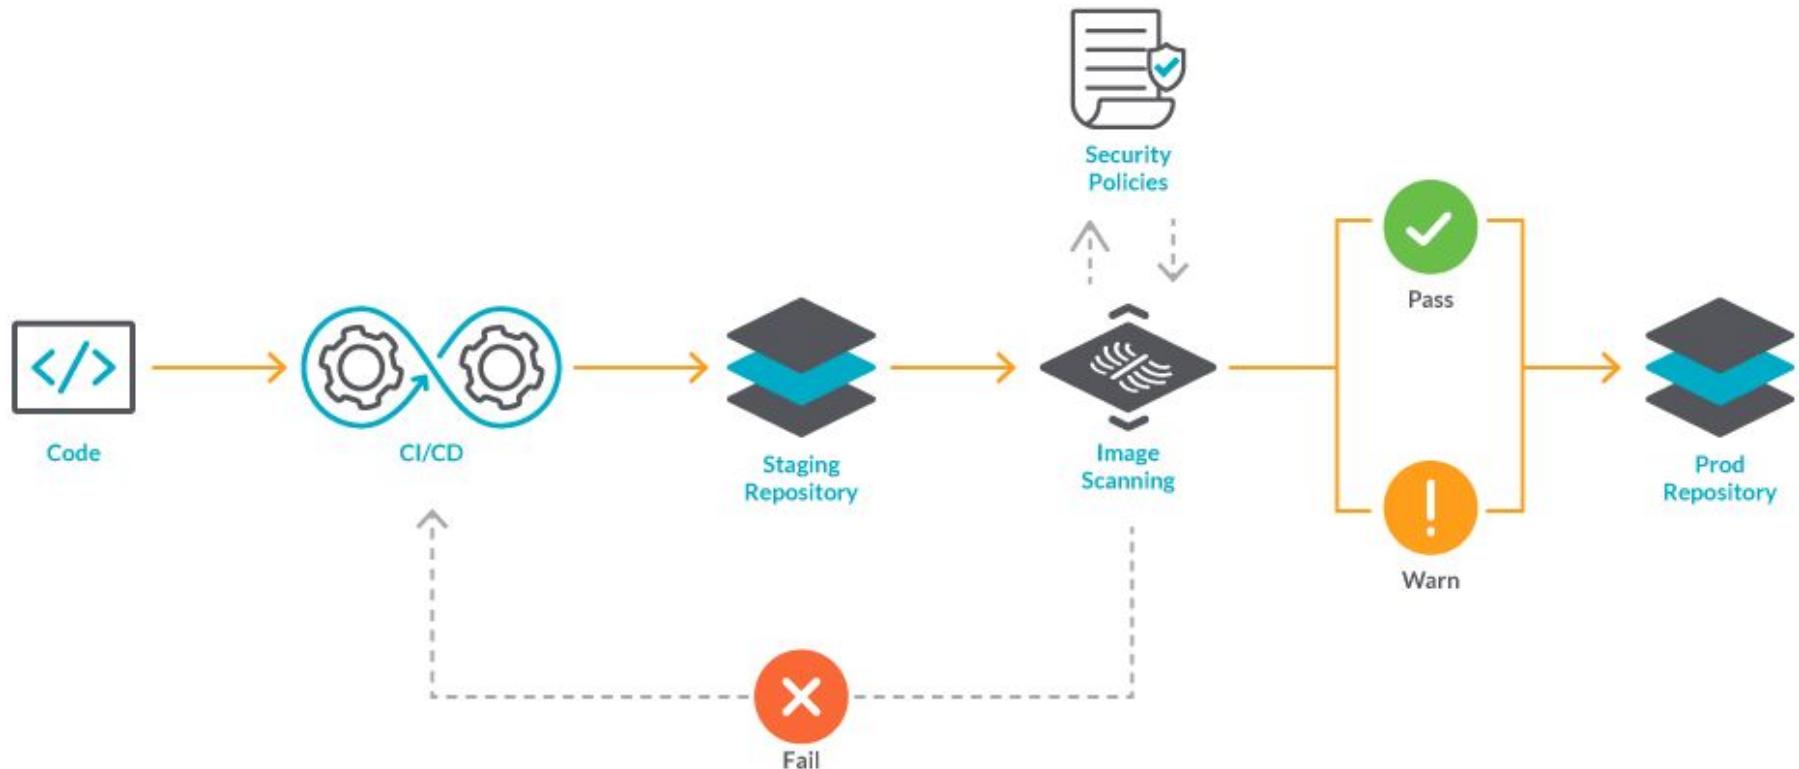
# SBOM (Software Bill of Materials)

Machine-readable list of:

- libraries
- dependencies
- versions
- Generated at build time

Audit value

- Transparency of software composition
- Faster incident response (e.g. Log4Shell)
- Regulatory compliance support



---

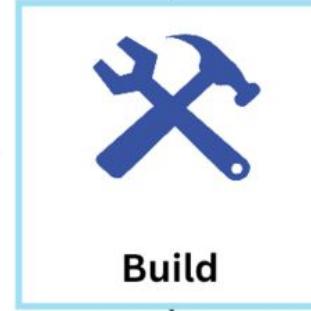
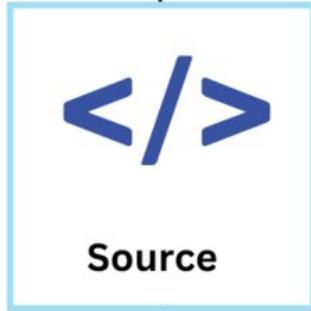
# Supply chain security

# What Is Supply Chain Security?

- Protects the entire path:
  - source code
  - build process
  - artifacts
  - deployment
- Assumes attackers may target CI/CD systems, not just apps



# Software Supply Chain Security



Managing access to IDEs

Source code review

Source code management (SCM)

CI/CD Pipeline

Containers & Registries

Dependencies

SBOMs

Code provenance & signature

Artifact repository

# SLSA (Supply-chain Levels for Software Artifacts)

- A framework defining maturity levels for build security
- Focuses on:
  - build provenance
  - reproducibility
  - tamper resistance

## Audit relevance

- Provides a measurable security maturity model
- Answers: How trustworthy is this artifact?

# in-toto

Framework for securing the software supply chain

- Verifies:
  - each step in the build pipeline
  - who performed it
  - what was produced

Audit value

- End-to-end traceability
- Strong non-repudiation of build steps

# Why Auditors Care

- Many modern attacks bypass runtime security
- Supply chain controls answer:
  - Who built this?
  - How was it built?
  - Has it been modified?

# Training Summary — What We Learned

- During this training, we learned:
- how cloud computing shifts security and control from manual operations to identity, policies, and APIs
- how cloud-native and microservice architectures, supported by 12-Factor App principles, improve isolation, resilience, and auditability
- how containers and Kubernetes manage applications through declarative, self-healing mechanisms
- how Deployments, probes, autoscaling (HPA/VPA), and rollouts enable controlled, auditable application lifecycle management
- how CI/CD pipelines and Infrastructure as Code turn change management into repeatable, reviewable, and enforceable processes
- how observability (logs, metrics, traces) provides evidence for incident analysis and compliance
- how Kubernetes RBAC and secret management enforce least privilege and protect sensitive data
- how image scanning, SBOMs, digital signatures, and supply-chain security frameworks (SLSA, in-toto) help ensure that only trusted software reaches production

# Final takeaway

Cloud-native security is not about adding more tools — it is about designing systems that generate control and audit evidence by default.